

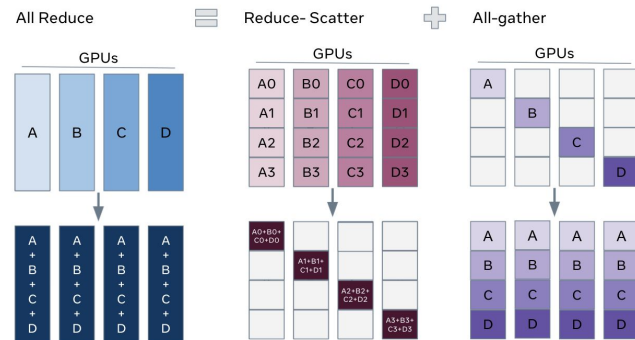
PyTorch Distributed and Parallel Training and FSDP

경희대학교
박현우

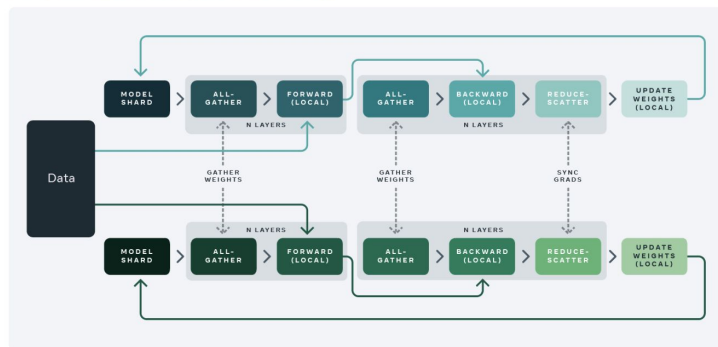
Table of Contents

Pytorch Distributed Training Basic

1. Overview of Pytorch Distributed Training
2. Collective Communications
3. Distributed Data Parallel (DDP)



Fully sharded data parallel training





FSDP

1. What is FSDP
2. How FSDP operates



1. Overview of Pytorch Distributed Training

PYTORCH

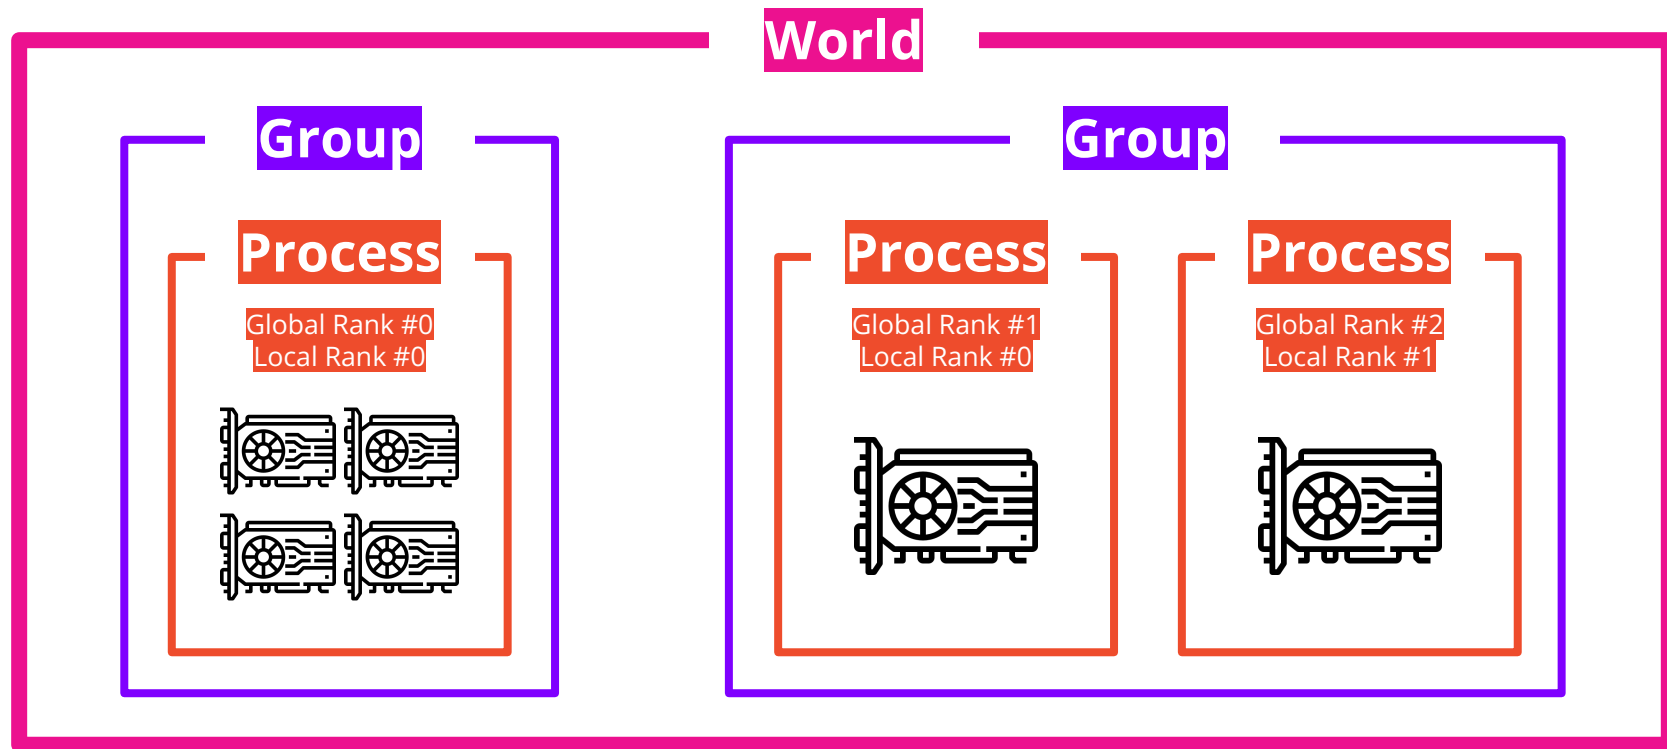


Terminology

Server/Node/Machine

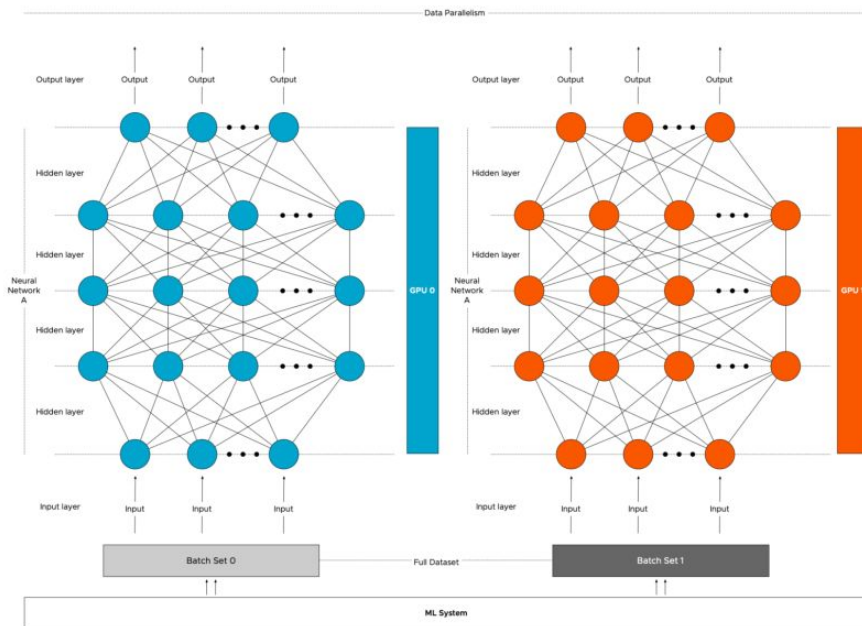


Terminology

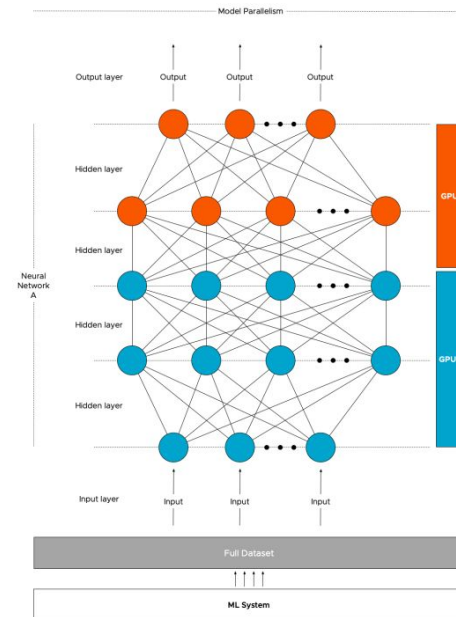


Terminology

Data Parallelism



Model Parallelism



1. torch.distributed

Low Level

High Level



Components

Modules

Paradigms

1.1 Main components of torch.distributed

DDP

Distributed
Data-Parallel
Training

RPC

Distributed
RPC
Framework

c10d

Communication
Core Library

1.1 Main components of torch.distributed

DDP

Aligned for **single-program multiple-data training** paradigm.

Specifically:

1. Take care of **gradient communication** to keep model replicas synchronized
2. **Overlap** communication with the gradient computations to speed up training.

1.1 Main components of torch.distributed

RPC

Aligned for supporting **general training structures** beyond data-parallel training.

For example,

1. Distributed Pipeline Parallelism
2. **Parameter Server Paradigm**
3. Combinations of DDP with other training paradigms




Fig. 2: DistDGL's logical components.

1.1 Main components of torch.distributed

RPC

Specifically:

1. Help manage remote object lifetime (RRef)
2. Extend the autograd engine to Distributed Autograd Engine.

Internally, RPC Backend of PyTorch relies on TensorPipe, which is an implementation of communications like NVLink or TCP, specially for PyTorch tensor.

1.1 Main components of torch.distributed

c10d

Aligned for **sending/receiving tensors** across processes within a group.


c10d offers both

1. **Collective Communication APIs**
(e.g., [all_reduce](#) and [all_gather](#))
2. **P2P communication APIs**
(e.g., [send](#) and [isend](#)).


1.1 Main components of torch.distributed

c10d

Collective Communication



P2P Communication



1.2 Basic Modules for Data-Parallel training

	Single Machine	Multiple Machine
Single GPU	Vanilla PyTorch	
Multiple GPU	DataParallel (DP) DistributedDataParallel (DDP)	DDP + torchrun (torch.distributed.elastic)

1.2 Basic Modules for Data-Parallel training

Vanilla PyTorch

For simplest,
small-sized model.

DataParallel (DP)

For speed up training with
minimal code changes.

1.2 Basic Modules for Data-Parallel training

DistributedDataParallel
(DDP)

For faster speed than DP,
but with more code.

DDP + torchrun
(torch.distributed.elastic)

For scaling across cluster of machines;
fine-grained error handling and dynamic
allocation/drop of machines during
training.

1.2 Basic Modules for Data-Parallel training

In case of DataParallel (DP)

- WARNING

It is recommended to use `DistributedDataParallel`, instead of this class, to do multi-GPU training, even if there is only a single node. See: [Use nn.parallel.DistributedDataParallel instead of multiprocessing or nn.DataParallel and Distributed Data Parallel](#).

1.3 Main Paradigms of Distributed/Parallel training

Data Parallel Trainings

DataParallel

Distributed
DataParallel

FSDP

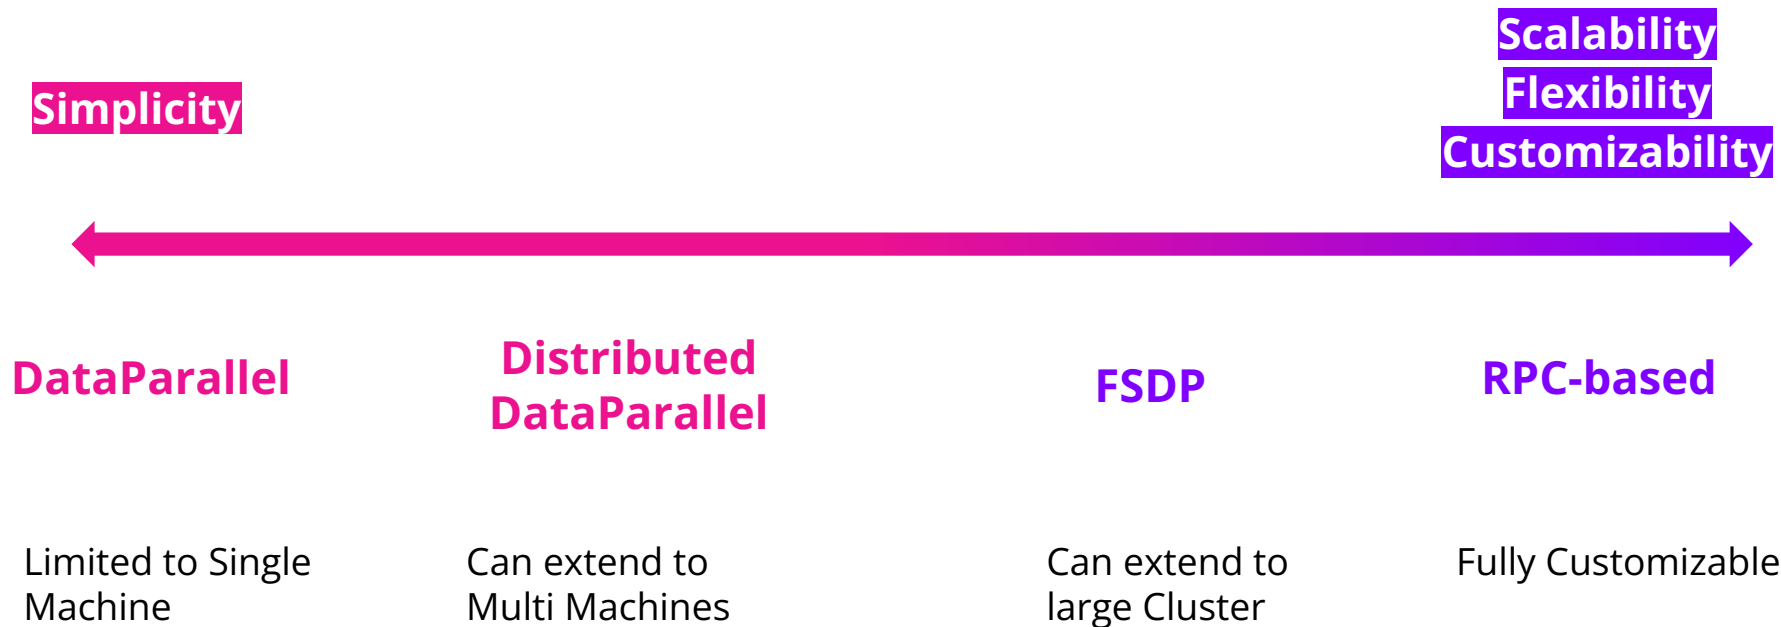
FSDP

RPC-based Trainings

Distributed
Pipeline Parallelism

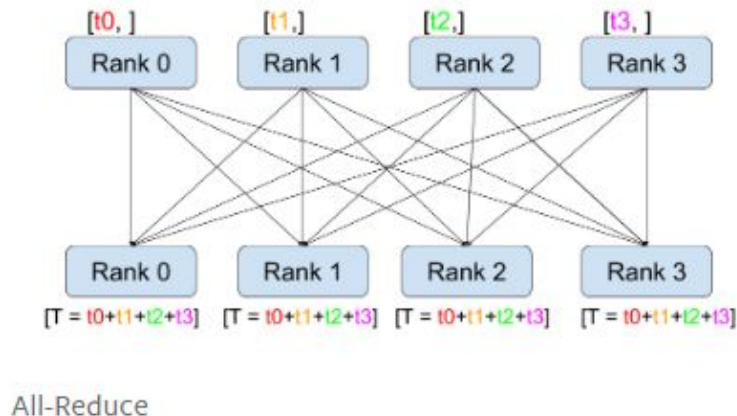
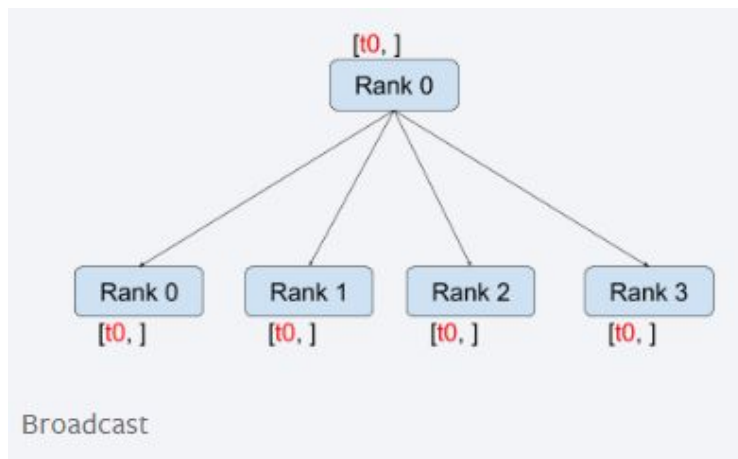
Parameter Server
Architecture

1.3 Main Paradigms of Distributed/Parallel training



2. Collective Communications

Collective Communications is to **communicate across multiple processes** in a cluster.



2. Collective Communications

Standard Communications

Scatter

Broadcast

All-Reduce

Gather

Reduce

All-Gather

Special Communications

Reduce-Scatter

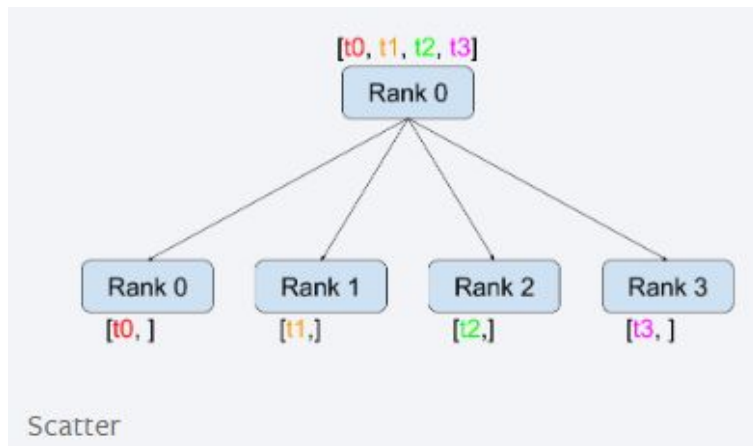
All-to-All

Barrier

2. Collective Communications

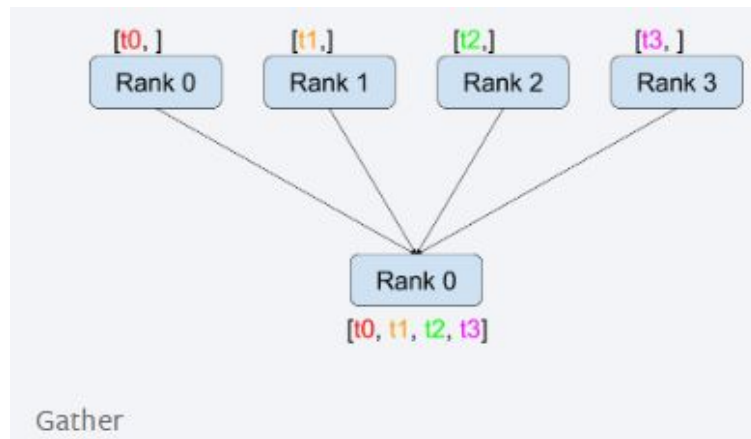
Scatter

Copies the i -th tensor to the i -th process.



Gather

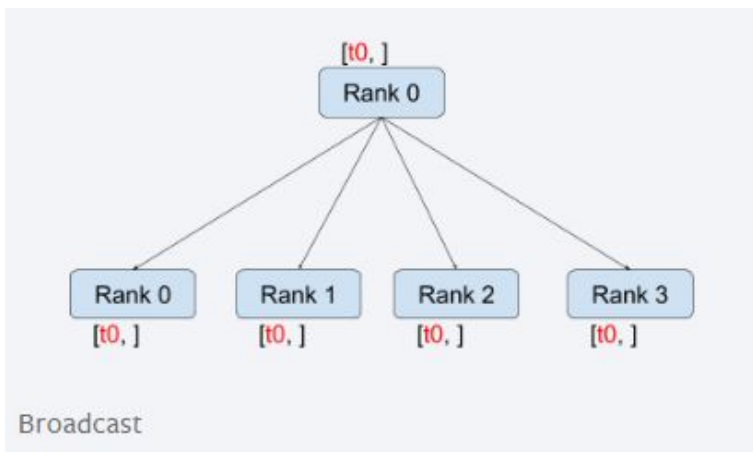
Copies tensor from all processes in dst.



2. Collective Communications


Broadcast

Copies tensor from source to all other processes.



Reduce

1. Applies operation to every tensor
2. Stores the result in destination.




Reduce

2. Collective Communications

All-Reduce


Same as Reduce, but the result is stored in all processes.



All-Reduce

All-Gather

Simply, All-Reduce without op.




All-Gather

2.1. All-Reduce Implementation


Approach #1

All processes are individually communicating to each other.



Approach #2

Aggregate on Master process
and propagate
(Reduce + Broadcast)



2.1. All-Reduce Implementation

Approach #1

Have $O(N^2)$
communications

Approach #2

Impose severe load on
master process

2.1. All-Reduce Implementation

As such, All-Reduce/All-Gather communications are **Extremely expensive**.

Thus, in order to efficiently distribute trainings,

**Need to Decompose these operations
and Distribute them in parallel.**

3. Distributed Data Parallel (DDP)

DDP implements both **Data and Model Parallelism** across multiple machines.

DDP can be run on in both

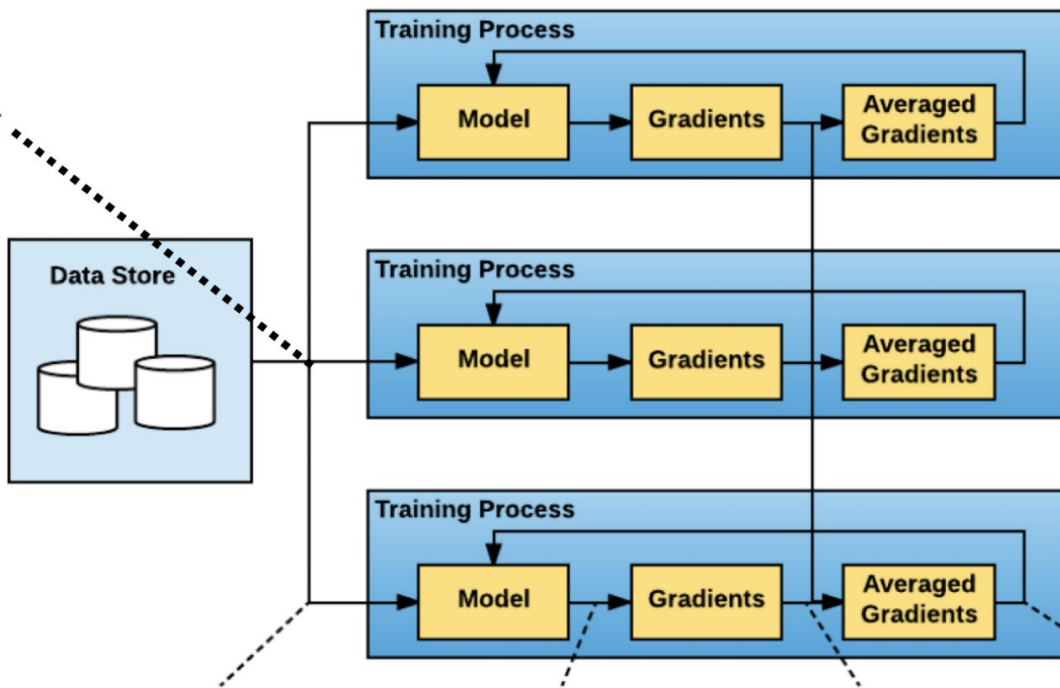
1. Single Machine - Multiple Devices
2. Multiple Machines - Multiple Devices

(in which for latter, **torchrun** is required.)

Also, DDP allows for **Heterogeneous cluster** (#).

3.1 Procedure of DDP

1. 데이터 Scatter



2. Forward 수행

3. Backward 수행

4. Gradient
All-reduce

5. 파라미터 업데이트

3.1 Procedure of DDP

Specifically:

1. DDP registers an **autograd hook** for each parameter (in `model.parameters()`)
2. The hook will fire when the corresponding gradient is computed in the **backward pass**.
3. Then DDP uses that signal to **trigger gradient synchronization** across processes.
4. Each process **updates** its own model with given (reduced) gradient.

3.2 Caveats of DDP

**Skewed
Processing
Speed problem**

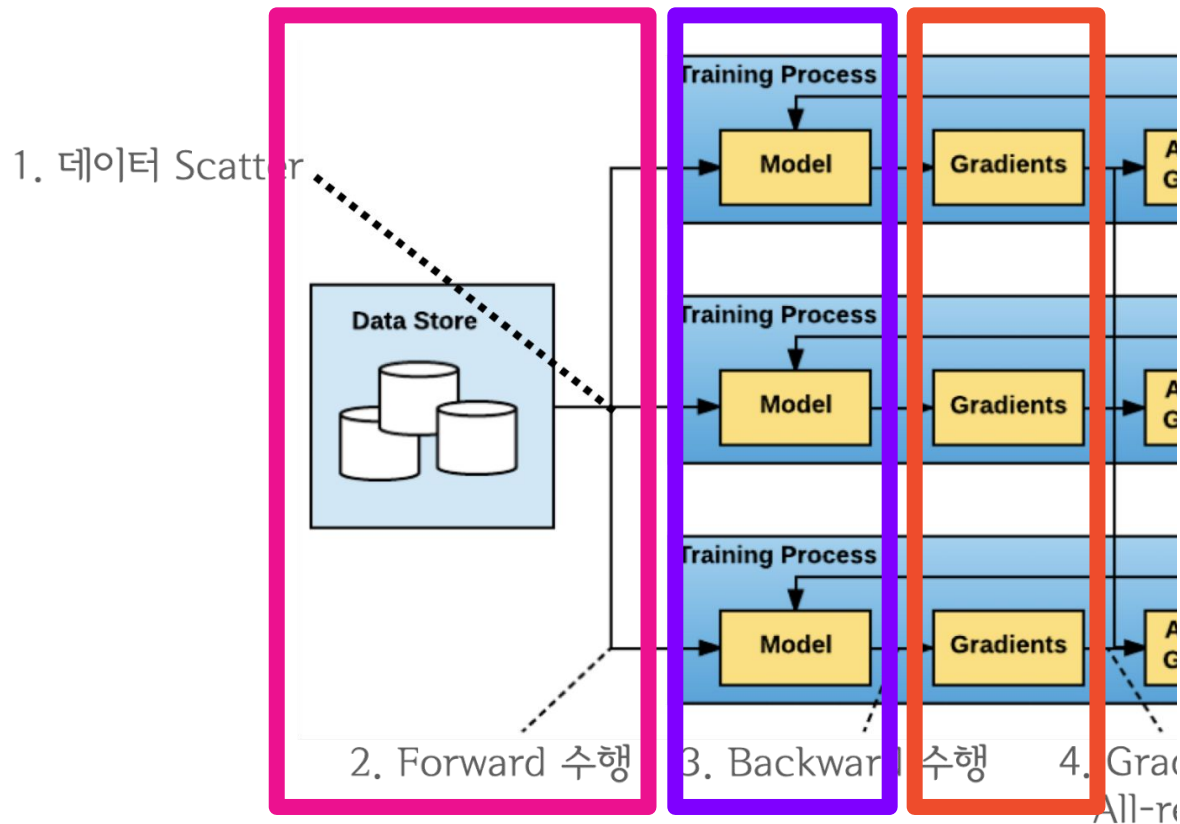
**Inefficient
Communication**

3.2.1 Skewed processing speed problem

There are 3 distributed synchronization points in DDP:

1. **Constructor**
2. **Forward pass**
3. **Backward pass**

3.2.1 Skewed processing speed problem



Constructor

Forward Pass

Backward Pass

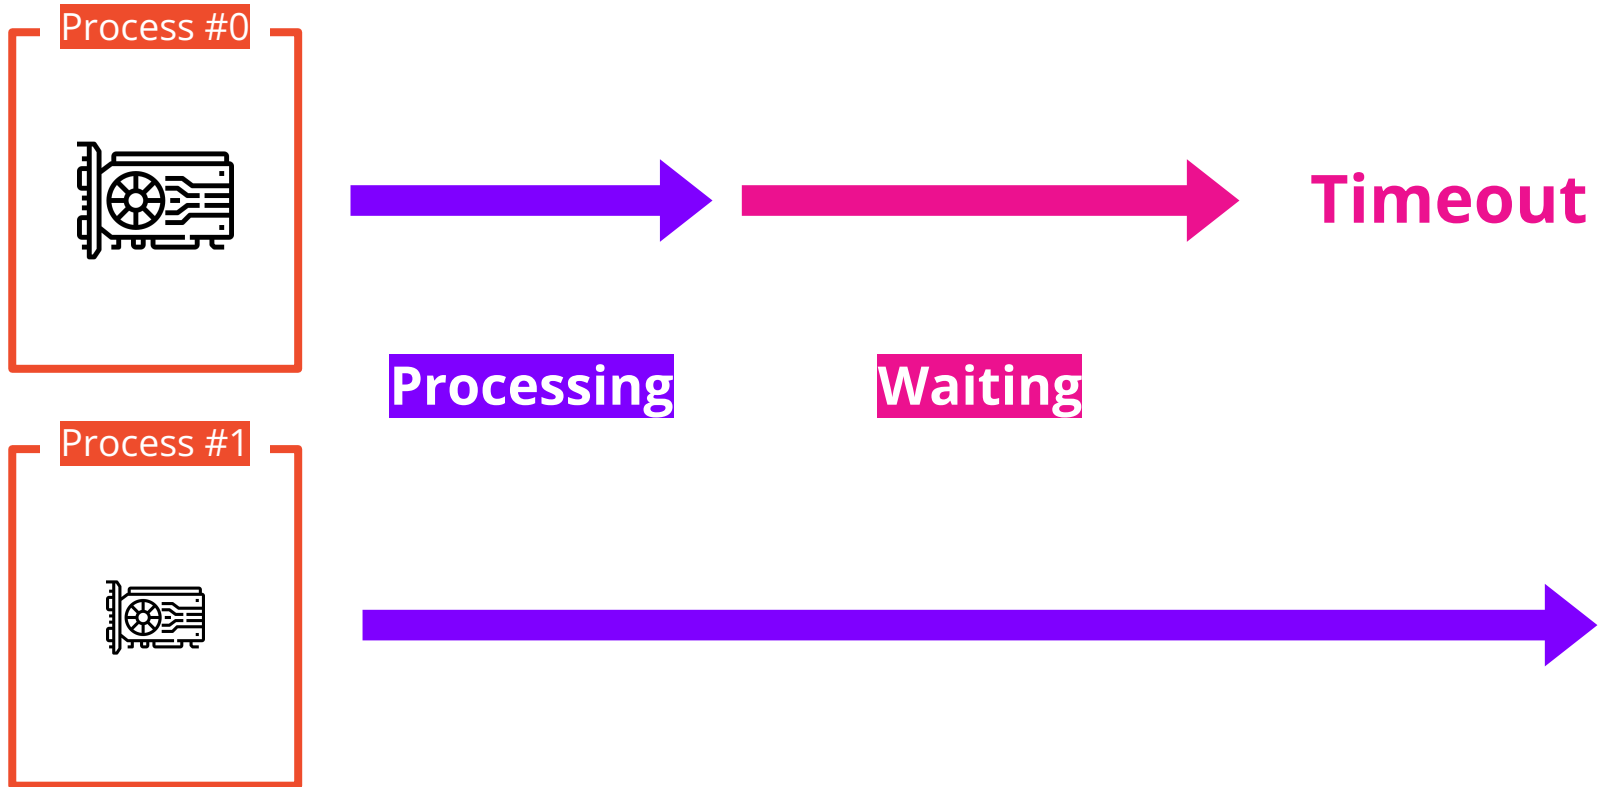
3.2.1 Skewed processing speed problem

In ideal, each process would

1. launch the same number of synchronizations
2. and reach these synchronization points in the same order.

**However, in practice,
Desynchronization can be occurred.**

3.2.1 Skewed processing speed problem



3.2.1 Skewed processing speed problem

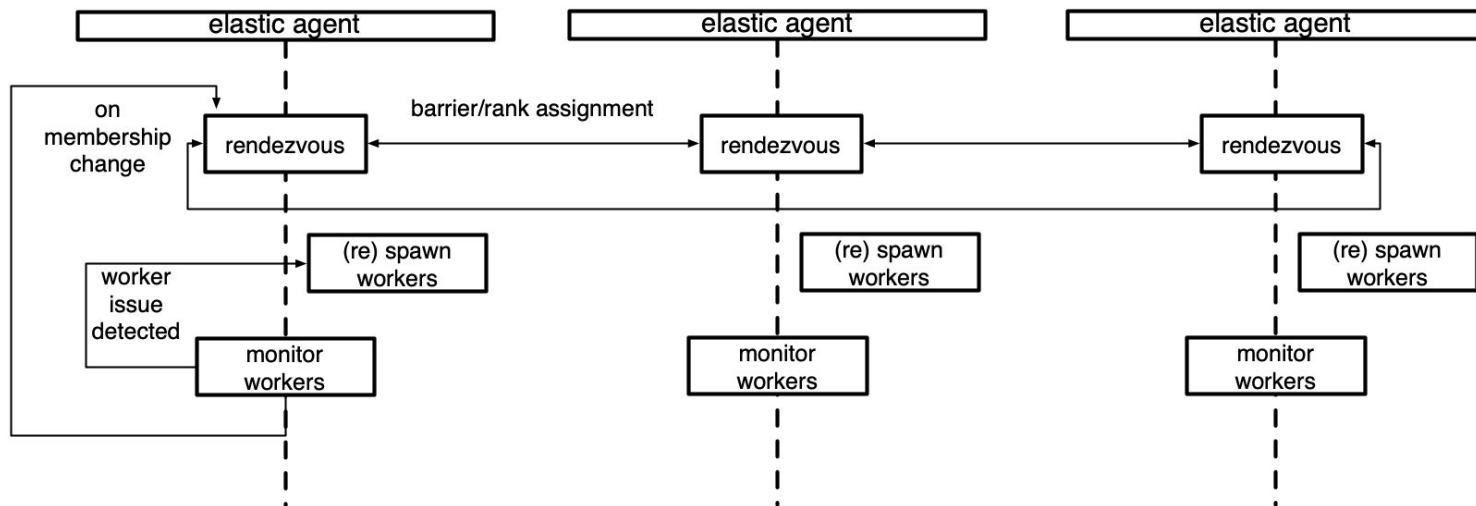
Such skewed processing speeds can also be occurred by

1. **Network delays**
2. **Resource contentions**
3. **Unpredictable workload spikes**

3.2.1 Skewed processing speed problem

DDP doesn't provide delicate, flexible **synchronization policy** across processes.

Unlike **torch.distributed.elastic**, DDP cannot recover from such timeout **by itself**.



3.2.1 Skewed processing speed problem

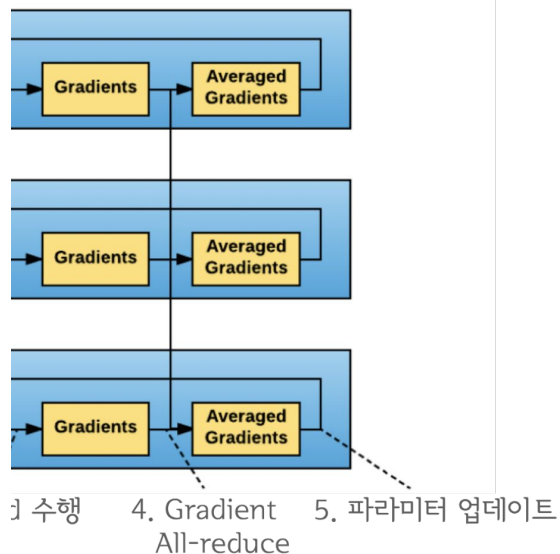
Thus, it is **fully responsible for developers** to deal with such synchronization.

• WARNING

This module assumes all parameters are registered in the model of each distributed processes are in the same order. The module itself will conduct gradient `allreduce` following the reverse order of the registered parameters of the model. In other words, it is users' responsibility to ensure that each distributed process has the exact same model and thus the exact same parameter registration order.

3.2.2 Inefficient Communication

DDP depends on **All-Reduce** to synchronize Updated gradients, which are **very expensive** operation.



3.3. Model Parallelism

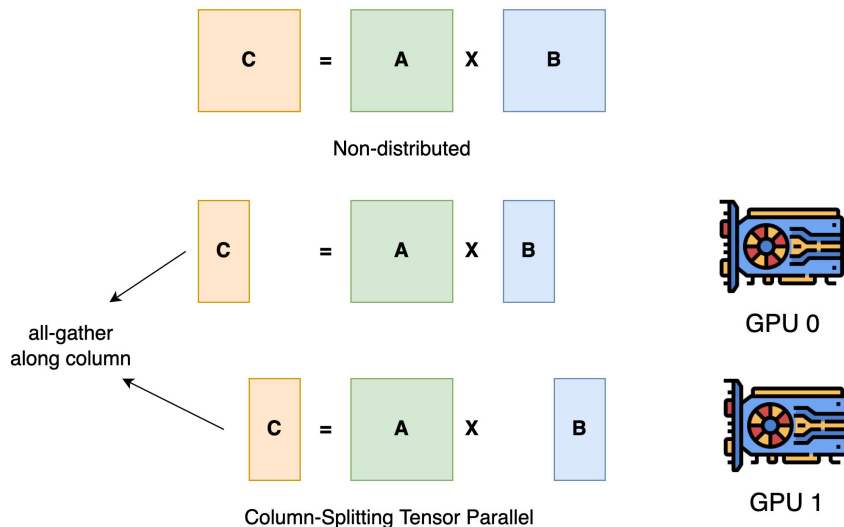
Types of Model Parallelism

**Intra-layer model
parallelism
or Tensor Parallelism**

**Inter-layer model
parallelism
or Pipeline Parallelism**

3.3. Model Parallelism in DDP

Intra-layer model parallelism




Vertically splitting the model

No Dependency among processes

3.3. Model Parallelism in DDP

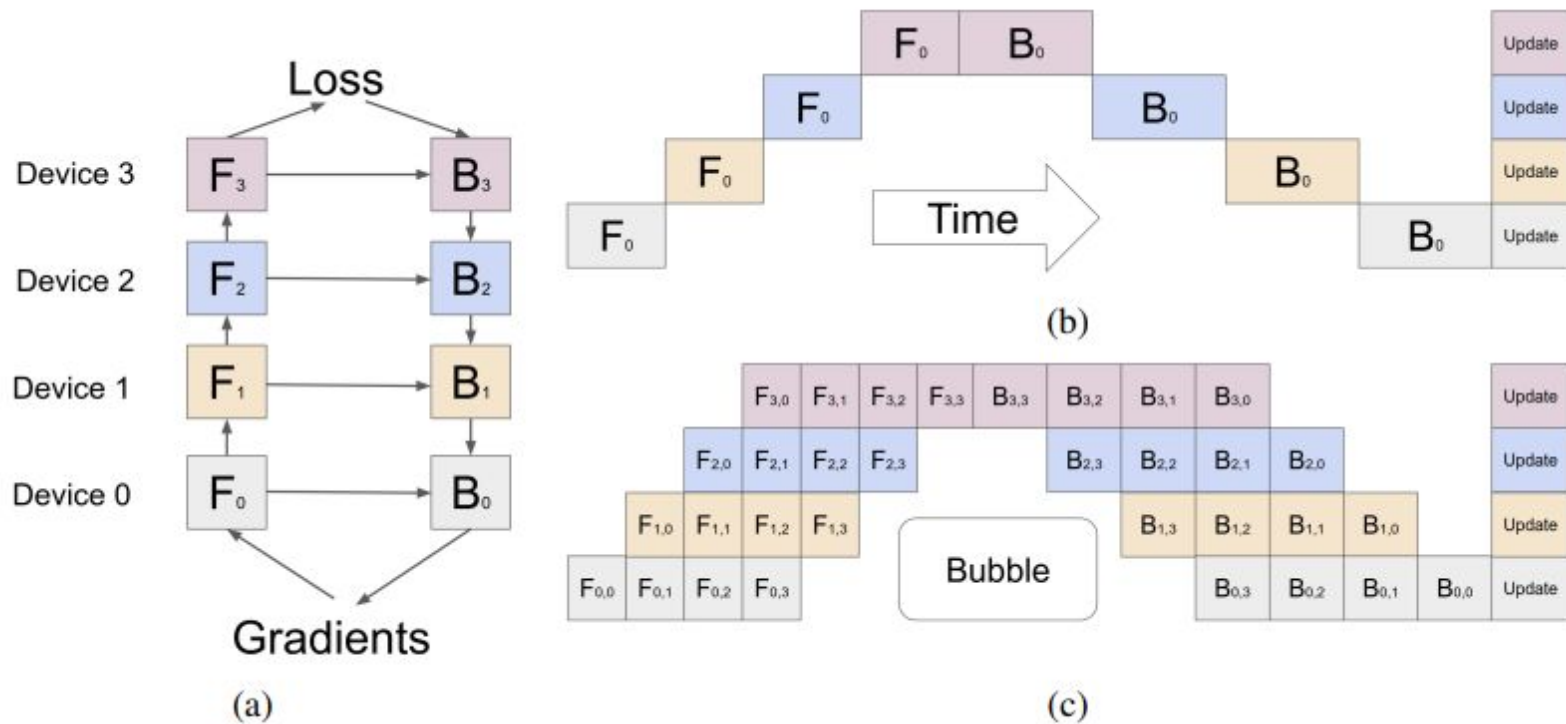
Inter-layer model parallelism




Horizontally splitting the model

Have Dependency among processes

3.3. Model Parallelism in DDP




3.3. Model Parallelism in DDP



**Inter-layer model
parallelism**

3.3. Model Parallelism in DDP



3.3. Model Parallelism in DDP

• NOTE

`Pipe` only supports intra-node pipelining currently, but will be expanded to support inter-node pipelining in the future. The forward function returns an `RRef` to allow for inter-node pipelining in the future, where the output might be on a remote host. For intra-node pipelining you can use `local_value()` to retrieve the output locally.

• NOTE

You can wrap a `Pipe` model with `torch.nn.parallel.DistributedDataParallel` only when the checkpoint parameter of `Pipe` is `'never'`.

3.3. Model Parallelism in DDP



Youhe-Jiang commented on Jun 30, 2022 • edited ▾

...

As shown in the figure below, pytorch does not support inter-node pipelining. I want to know if fairscale supports it? Thanks for any replies!

```
.. note::  
    :class:`Pipe` only supports intra-node pipelining currently, but  
    will be expanded to support inter-node pipelining in the future.  
    The forward function returns an :class:`~torch.distributed.rpc.RRef`  
    to allow for inter-node pipelining in the future, where the output  
    might be on a remote host. For intra-node pipelining you can use  
    :meth:`~torch.distributed.rpc.RRef.local_value` to retrieve the  
    output locally.
```



min-xu-ai commented on Jul 1, 2022

Contributor ...

There is some experimental support for multiprocessing pipe. Unfortunately, we don't have anyone who is familiar with it at the moment.


4. FSDP

Fully Sharded Data Parallel FSDP is Data and Model parallel training algorithm.


FSDP aims to train very large models on large cluster of machines.

FSDP solves problems in DDP, such as insufficient model parallelism support or inefficient communications.

4. FSDP



4.1. Full Parameter Sharding




Gradient

Gradient through either forward and backward pass of specific subset of the model (here, layers).

e.g.

1. Forward pass through only **Layer #1**.
2. Backward pass through **Layer #2** (given Gradient of **Layer #3**)

4.1. Full Parameter Sharding



Optimizer States

Gradients through backward pass of complete model.

i.e. Backward pass through all Layer #1, #2, #3.

= All backward Gradients

4.1. Full Parameter Sharding

Target of Replication

DDP

Full Model Weights

Full Optimizer States

FSDP

Partial Model Weights

Partial Optimizer States

Partial Gradients

4.1. Full Parameter Sharding

Full Parameter Sharding **shards across processes gradients of the model.**

The name of Full Parameter Sharding comes from that parameters are “fully sharded” across the cluster.

c.f. Replication of full parameters in DDP

Thus, FSDP can be regarded as native **Inter-layer Parallelism** or **Pipeline Parallelism**.

4.2. FSDP optimization strategy

FSDP also ensures that the computation for **each microbatch still local** to each GPU worker, even with Full Parameter Sharding. (#)

FSDP does so by **sharing parameters in every forward/backward compute**;

i.e. instead of sharing each microbatch.

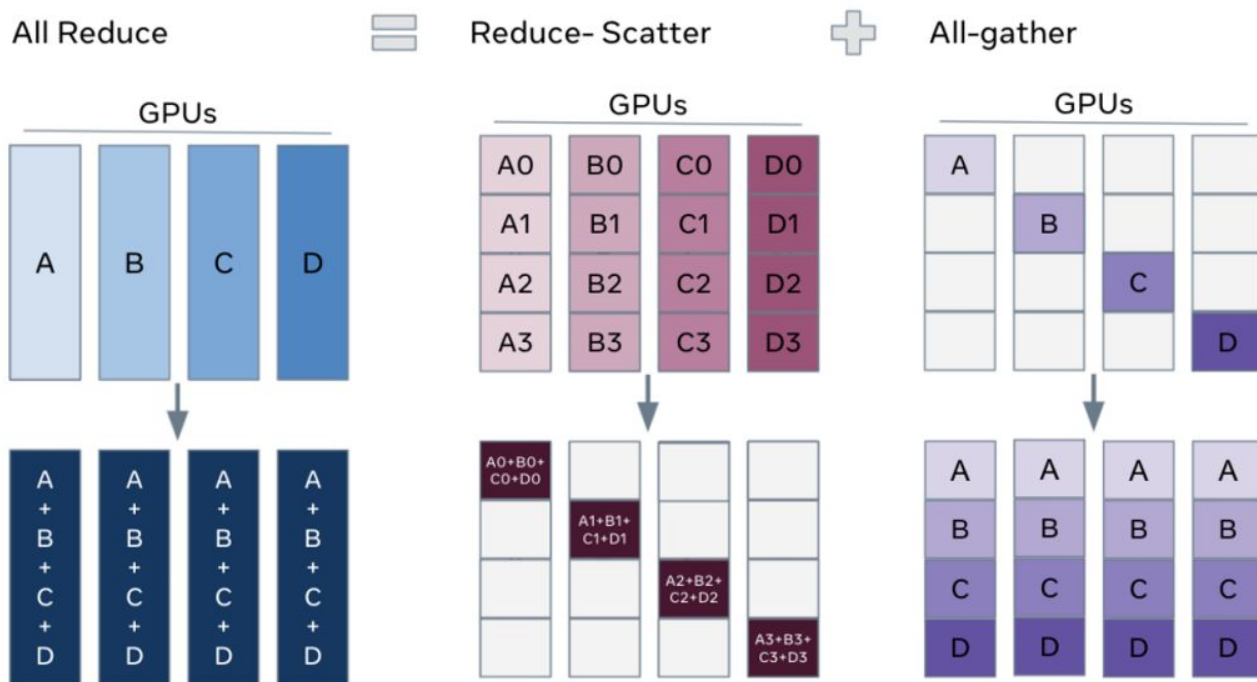
4.2. FSDP optimization strategy

Because FSDP synchronizes for Gradient, FSDP involves **larger communication volume** compared to DDP.

FSDP reduces such increased communication overhead by


**Decomposing
communication/computation
and Overlapping them in training**

5. How FSDP operates



5.1. Decomposition of All-Reduce

All Reduce



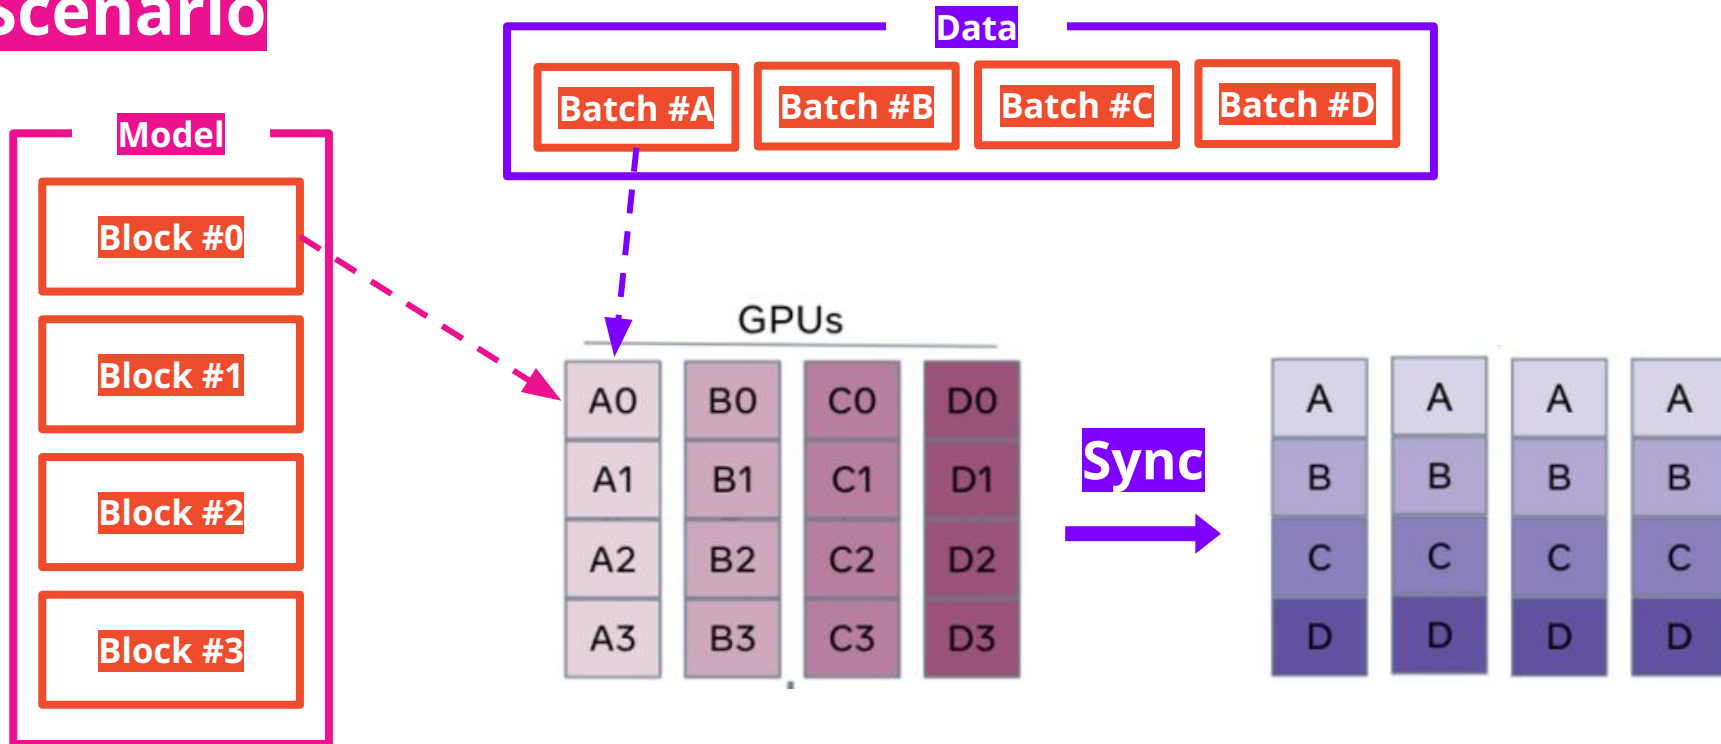
In standard DDP,

optimizer states with respect to each batches are sync-ed across all devices via All-Reduce.

In FSDP,
such optimizer states are decomposed
into multiple Gradients.

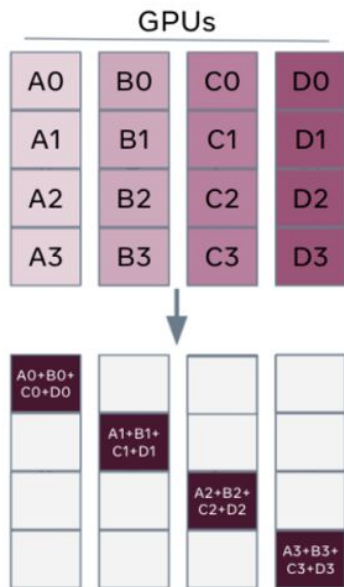
5.1. Decomposition of All-Reduce

Scenario



5.1. Decomposition of All-Reduce

Reduce- Scatter




During the Reduce-Scatter phase,

the gradients are aggregated in equal blocks among ranks based on their rank index.

5.1. Decomposition of All-Reduce

All-gather



During the all-gather phase,



the sharded portion of aggregated gradients on each GPU are propagated to all GPUs.

5.1. Decomposition of All-Reduce

Reduce- Scatter



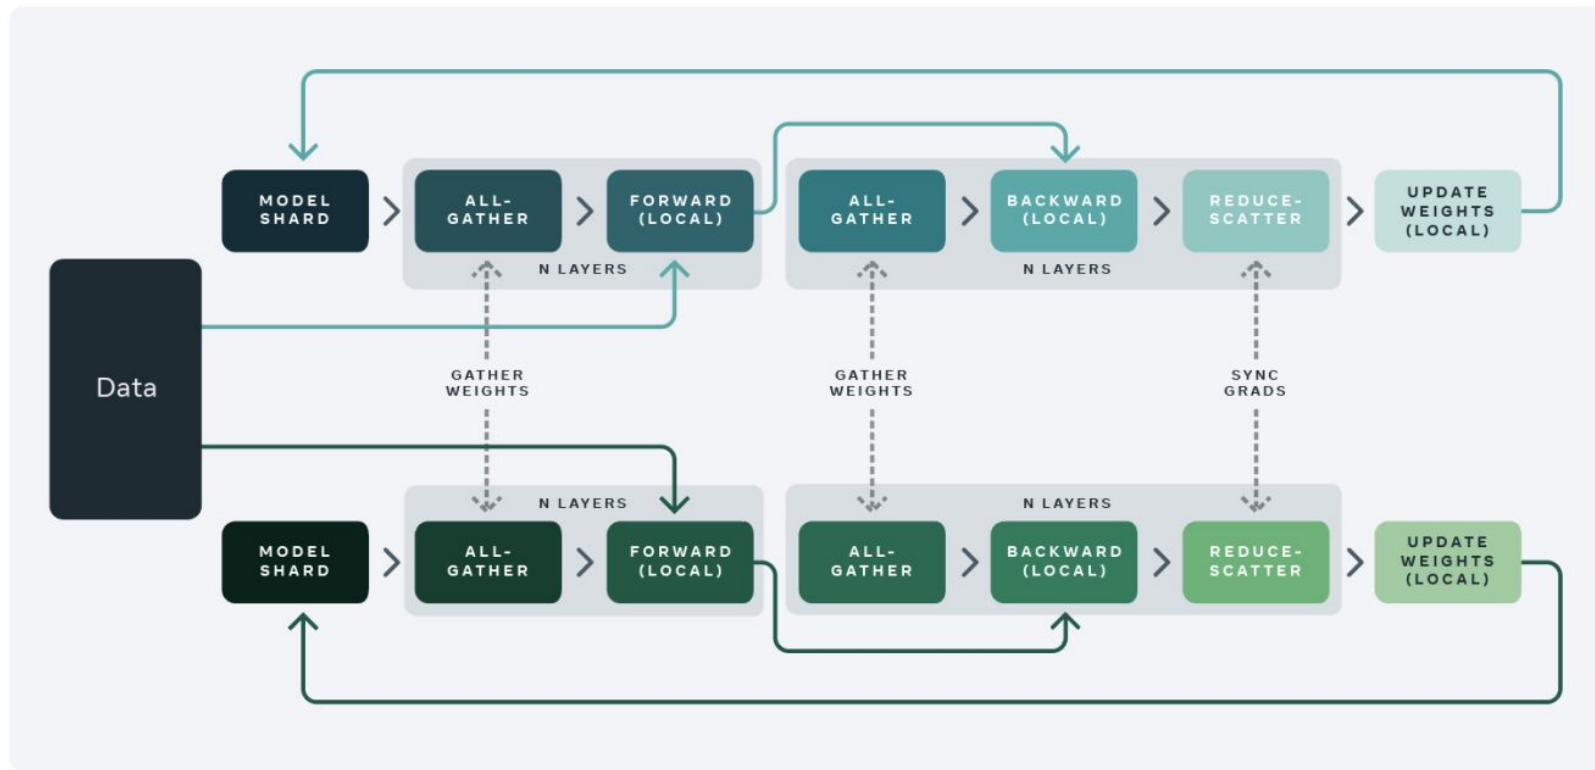
All-gather



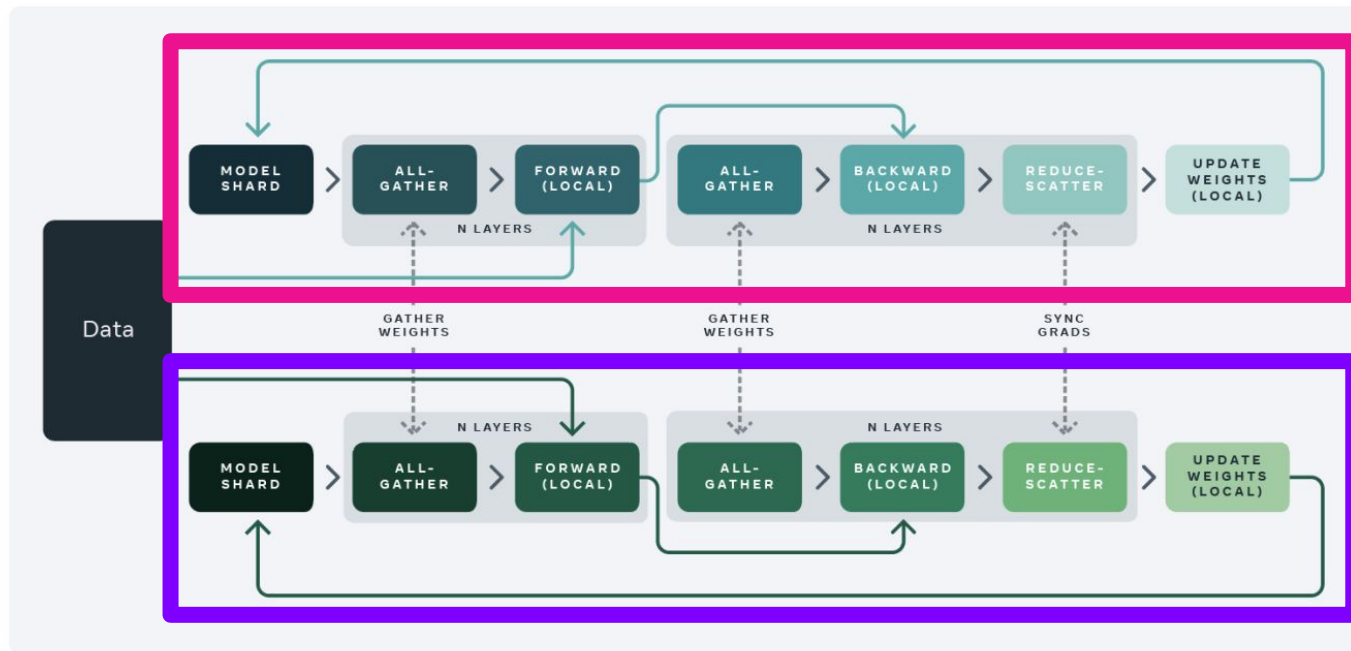
Once All-Reduce is decomposed into multiple Reduce-Scatter + All-Gather,

we can now Rearrange them and Run Asynchronously.

5.2. Structure of FSDP



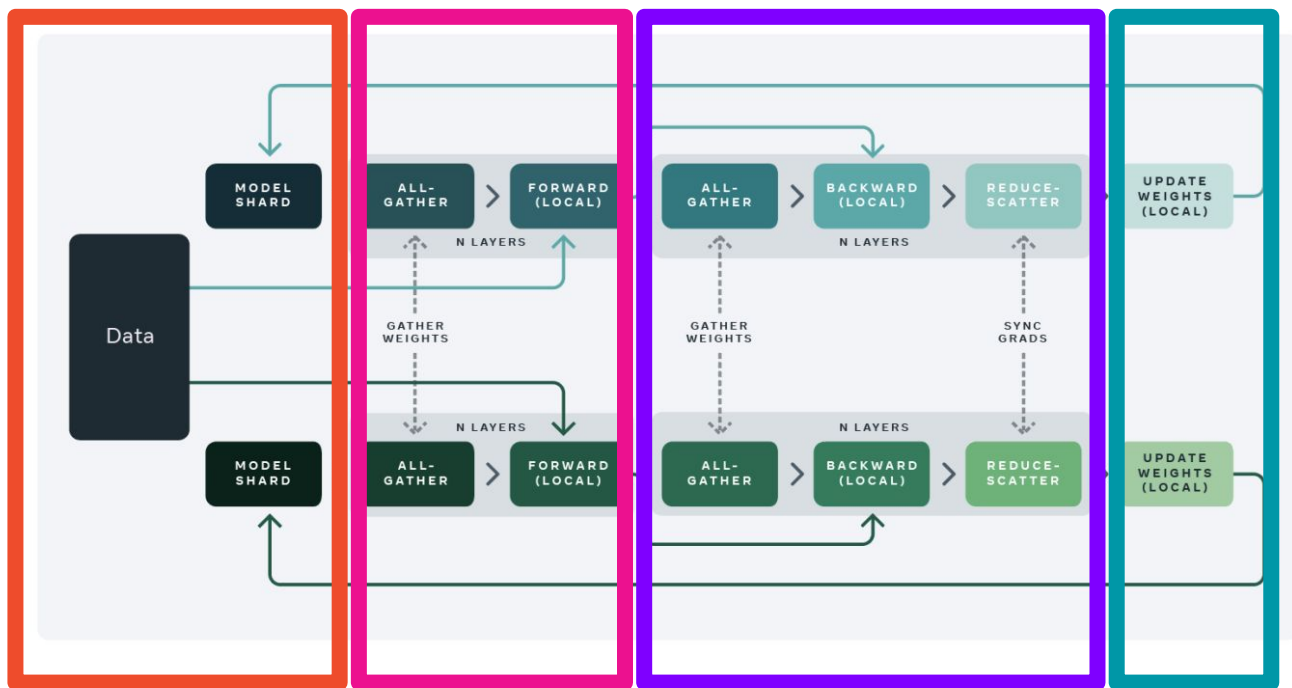
5.2. Structure of FSDP



Process #0

Process #1

5.2. Structure of FSDP



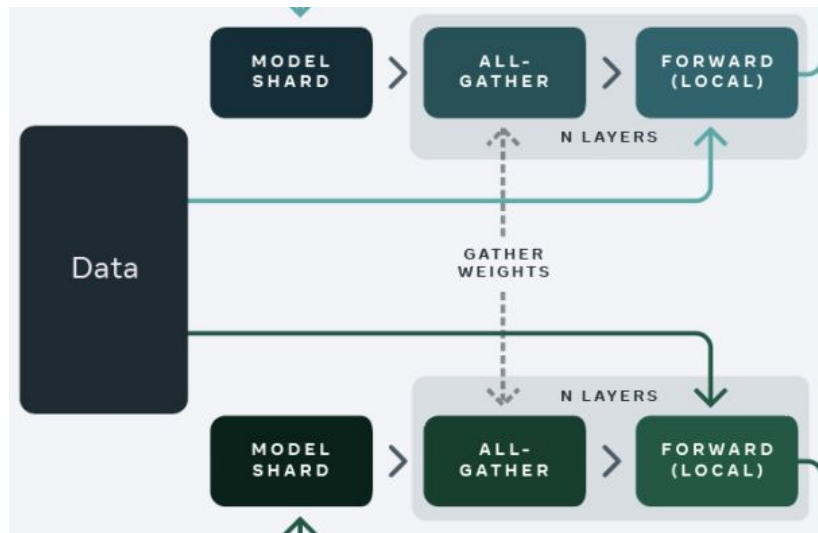
1. Construct

2. Forward Pass

3. Backward Pass

4. Update

5.2.1. Construct




Each process contains one block of the model;

e.g. Process #0 contains **Block #0**.

Dataset is also sharded for each process;

e.g. Process #0 contains **Batch #A**.

5.2.2. Forward Pass




In Forward Pass stage,

each process perform following steps for all layers of the model.

1. **Get layer using All-Gather**
2. **Compute forward pass**
3. **Discard gathered layer**

(Assume there are N layers in each block)

5.2.2. Forward Pass




Step 1

For Layer #0 (in **Block #0**),
Process #1 calls **All-Gather for Layer #0** as Process #1
doesn't have it.

(Process #0 can skip this, as Process #0 holds it on
itself.)

5.2.2. Forward Pass




Step 2

All Processes compute for Layer #0 on their own dataset.

e.g. Process #0 compute on **Batch #A**

5.2.2. Forward Pass




Step 3

Process #1 discards gathered layer, Layer #0.

This is to maximize memory efficiency.


5.2.2. Forward Pass



Pseudo Code

```
FSDP forward pass:  
  for layer_i in layers:  
    all-gather full weights for layer_i  
    forward pass for layer_i  
    discard full weights for layer_i
```

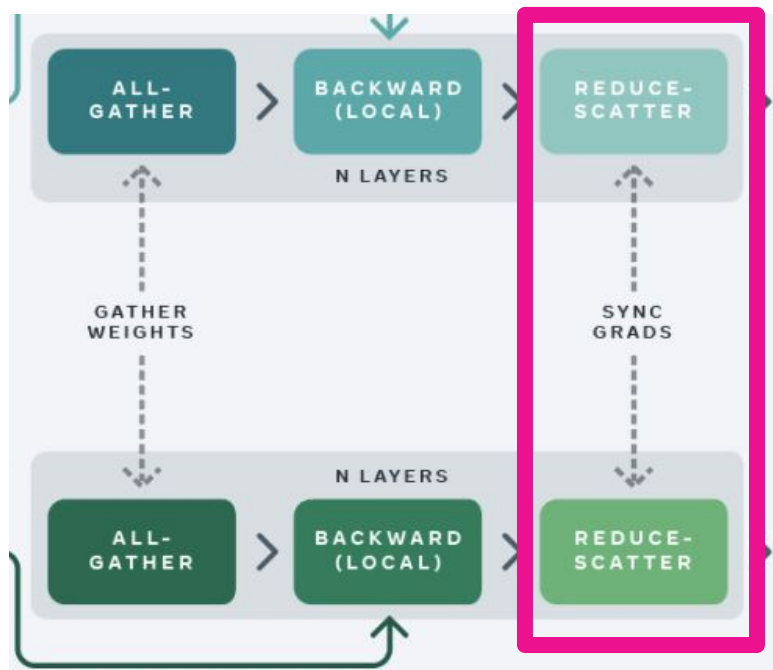
5.2.2. Forward Pass



At the end of forward pass,

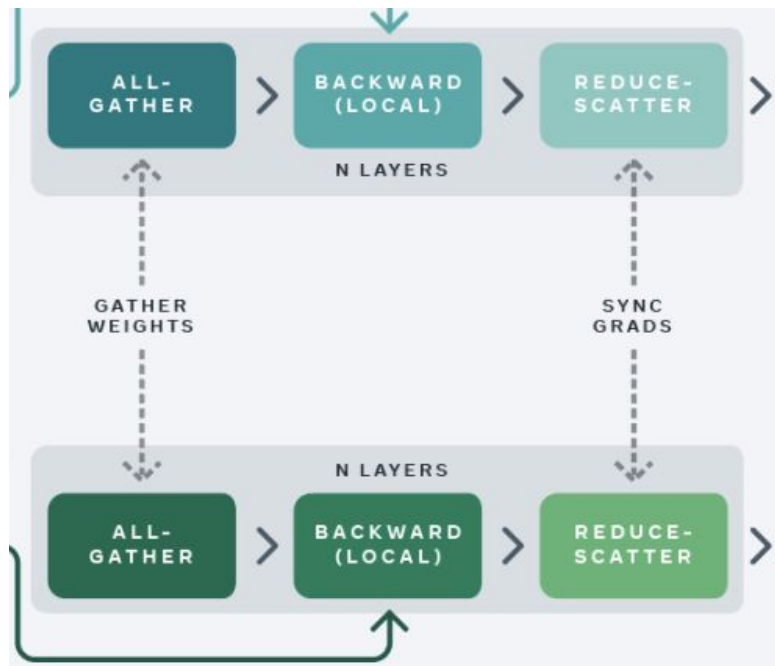
each process contains **outputs of full model** with respect to their own dataset.

5.2.3. Backward Pass



Computation of backward pass is the same as forward pass, except **Gradient Synchronization**.

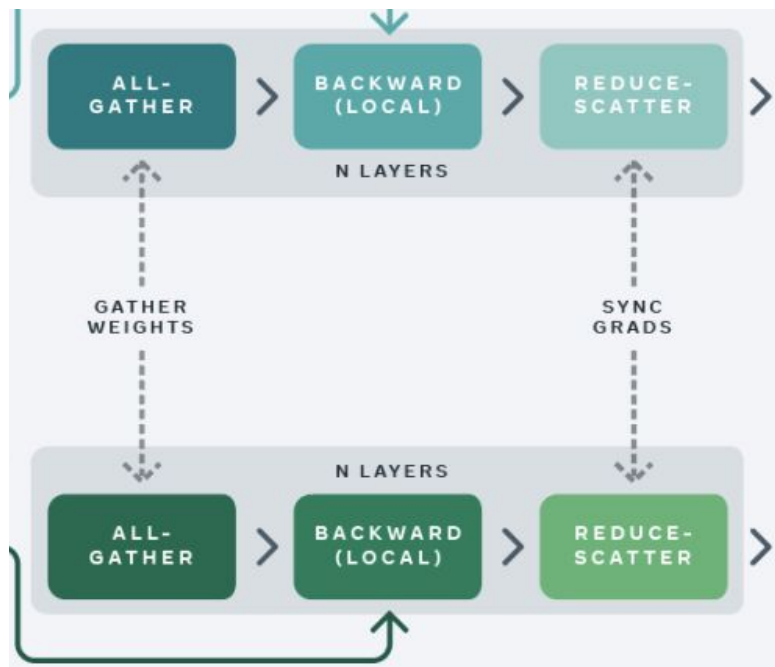
5.2.3. Backward Pass



Gradient Synchronization

After backward pass for each layer, the **local gradients are synchronized** across the processes via **reduce-scatter**.

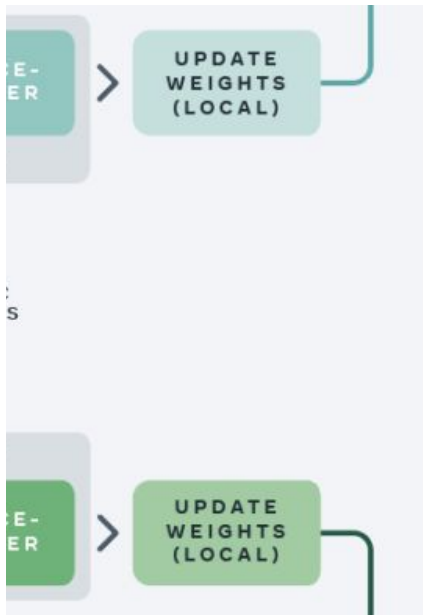
5.2.3. Backward Pass



Pseudo Code

```
FSDP backward pass:  
  for layer_i in layers:  
    all-gather full weights for layer_i  
    backward pass for layer_i  
    discard full weights for layer_i  
    reduce-scatter gradients for layer_i
```

5.2.4. Update



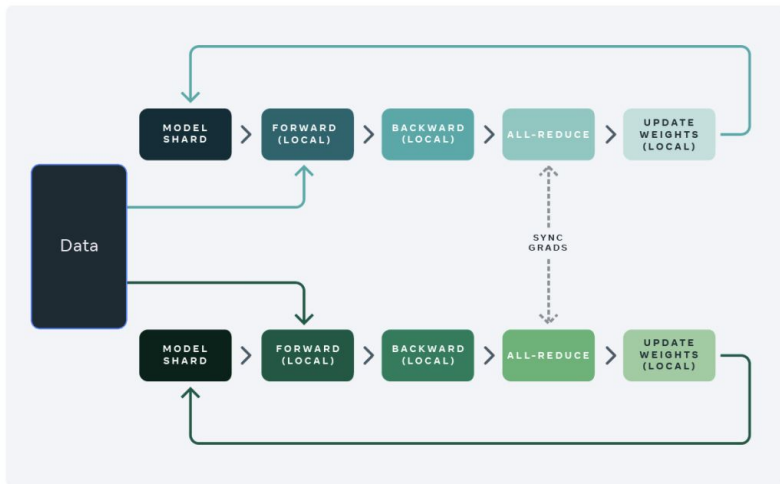
Weight Update

Given reduced gradients from other processes,

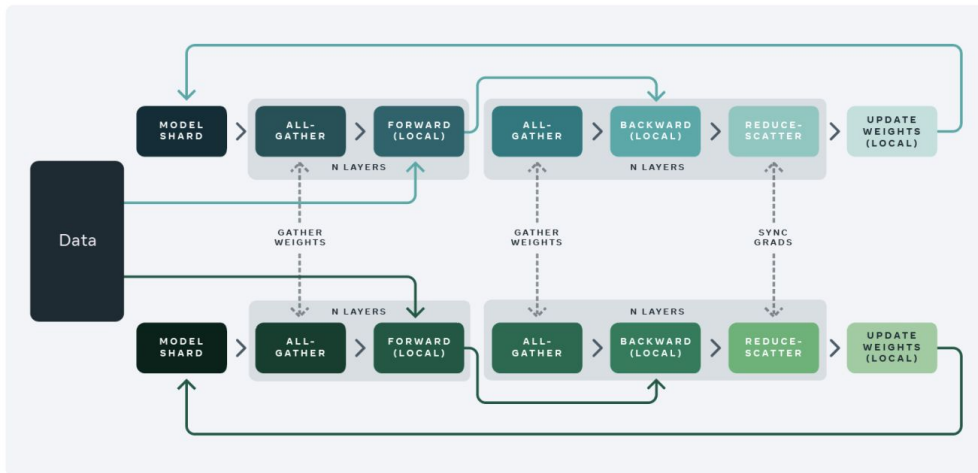
each process **updates its local weight.**

5.3. Comparison to DDP

Standard data parallel training

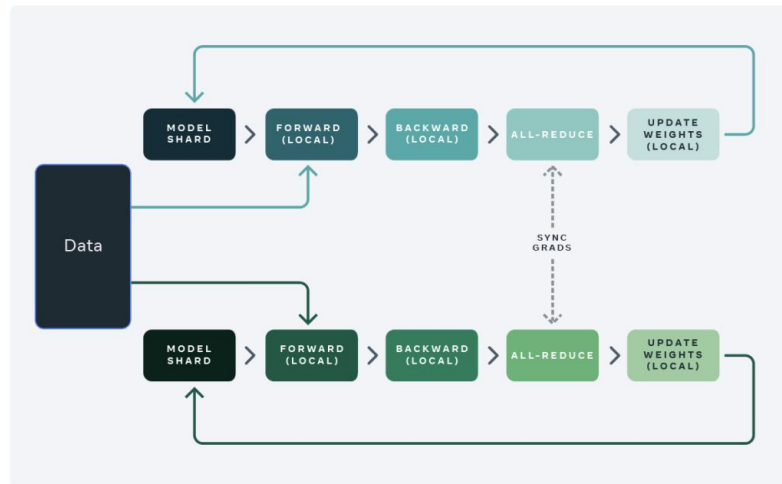


Fully shared data parallel training

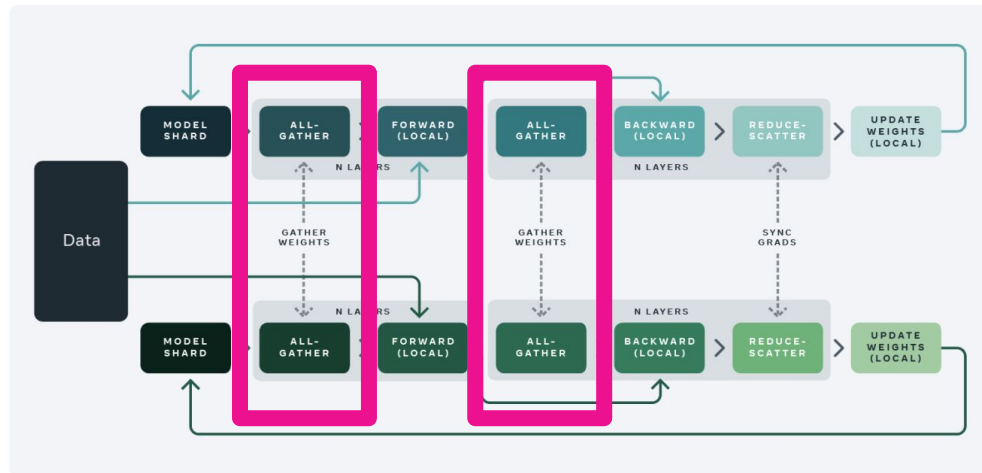


5.3. Comparison to DDP

Standard data parallel training



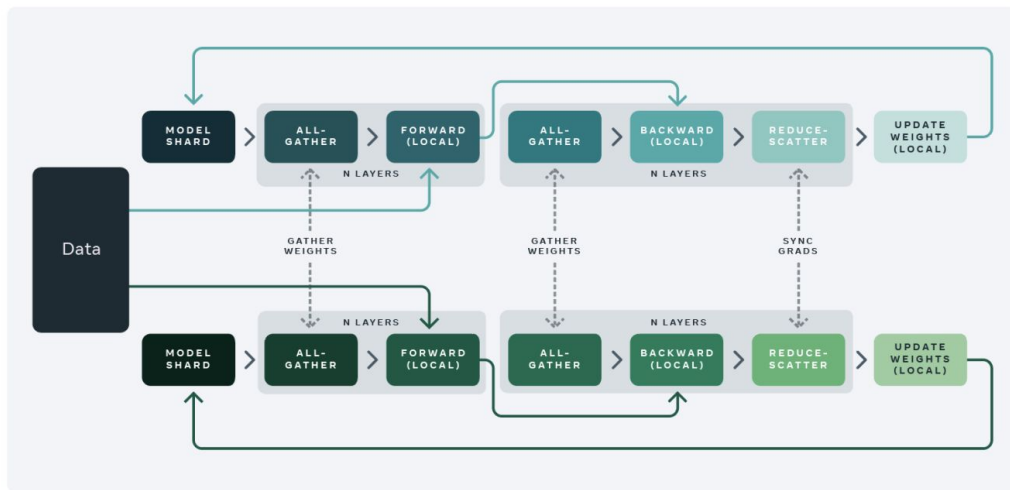
Fully sharded data parallel training



FSDP does have more communications than DDP.

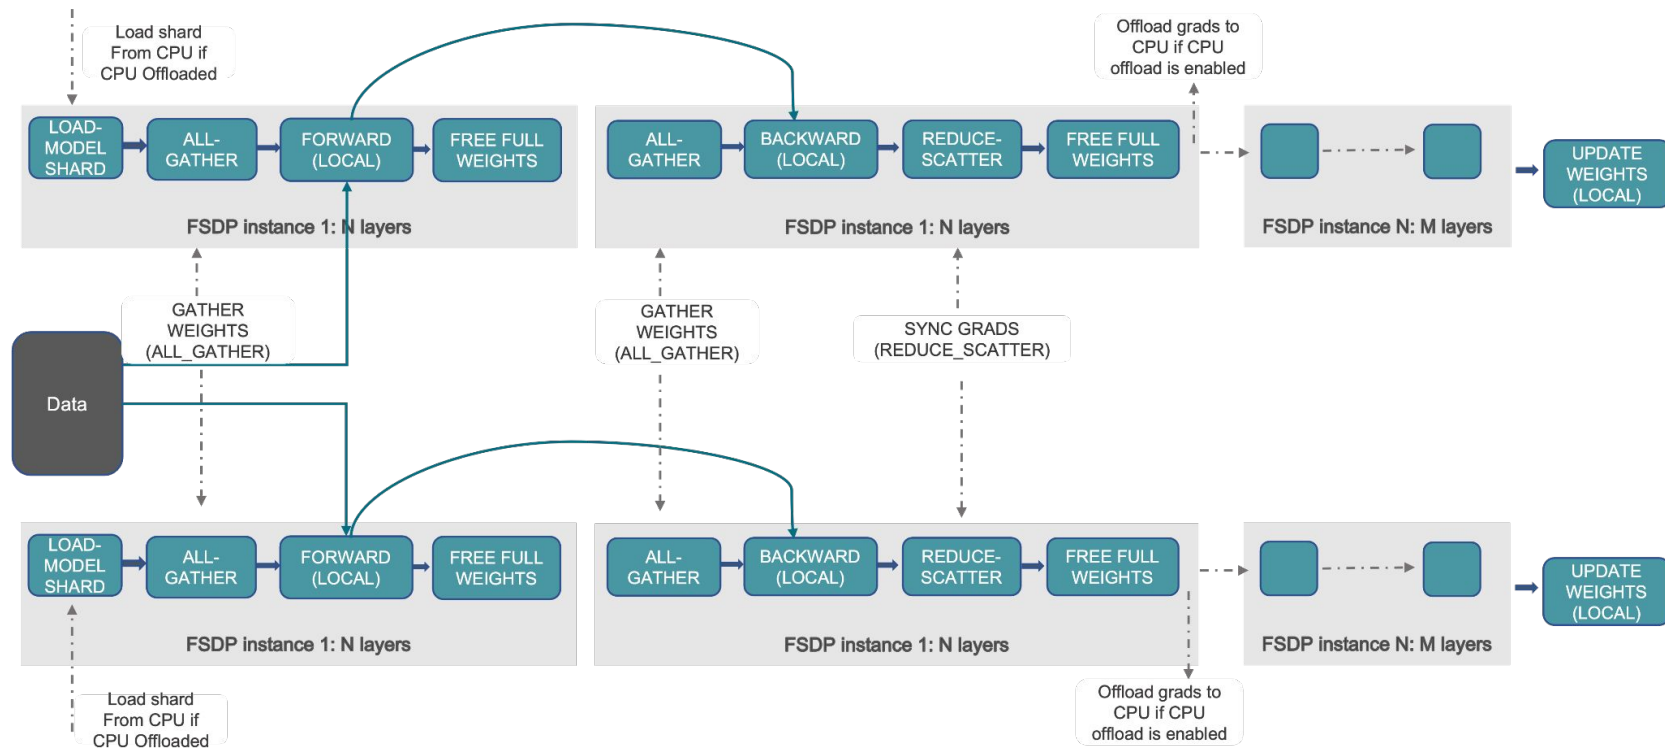
5.3. Comparison to DDP

Fully sharded data parallel training




However, FSDP can overlap communication with computation.

5.4. Full Architecture of FSDP




5.4. Full Architecture of FSDP



Although previous example has 1:N layers per FSDP instance,

FSDP can also **map N:M layers**.

5.4. Full Architecture of FSDP



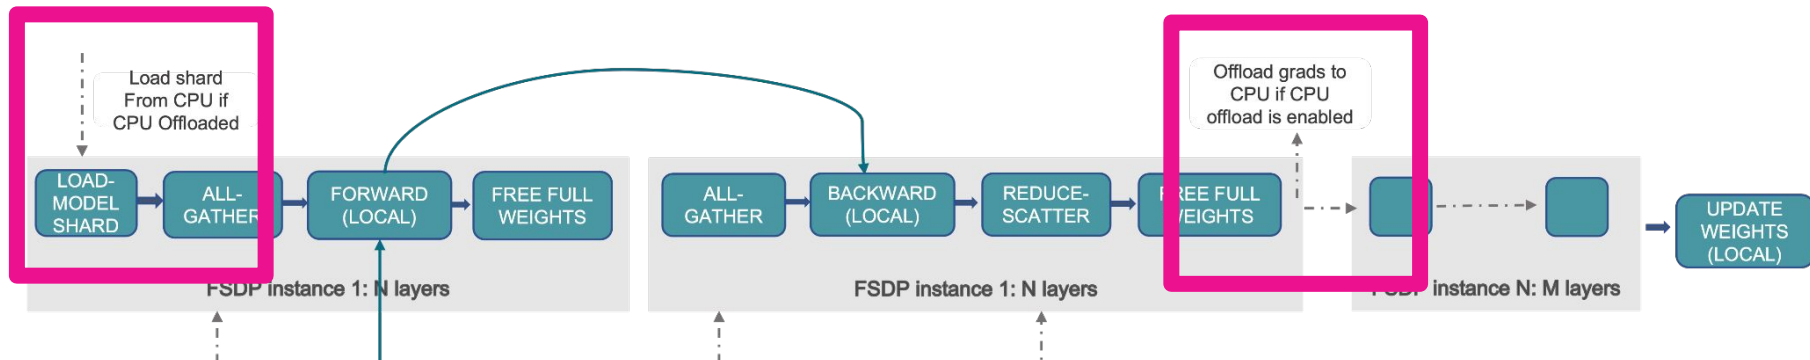
In this case,

although some instances contain redundant layers, (= **waste of memory**), they are cooperating in communication (= **share communication load**).

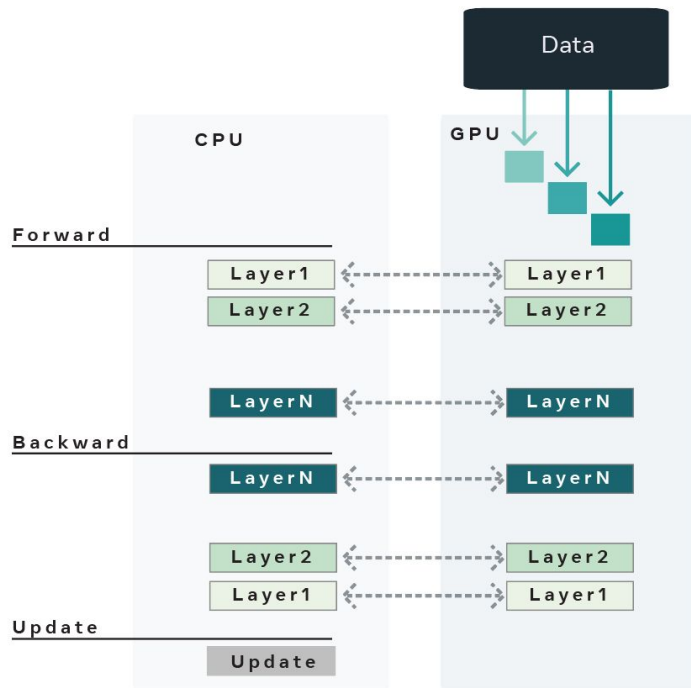
5.4. Full Architecture of FSDP

FSDP also supports **CPU-offload** (Zero-Offload)

in cases where **even sharded parameters are too large** to fit in each device.



5.4. Full Architecture of FSDP



CPU Offload for shard loading

CPU memory temporally contains model shard and sends copy to device.

5.4. Full Architecture of FSDP




Figure 3: ZeRO-Offload training process on a single GPU.

CPU Offload for Optimizer

CPU is responsible for updating the parameters and holding onto the optimizer state.

5.5. Performance Analysis




Figure 3. GPTSmall (125M) Throughput vs Number of GPUs

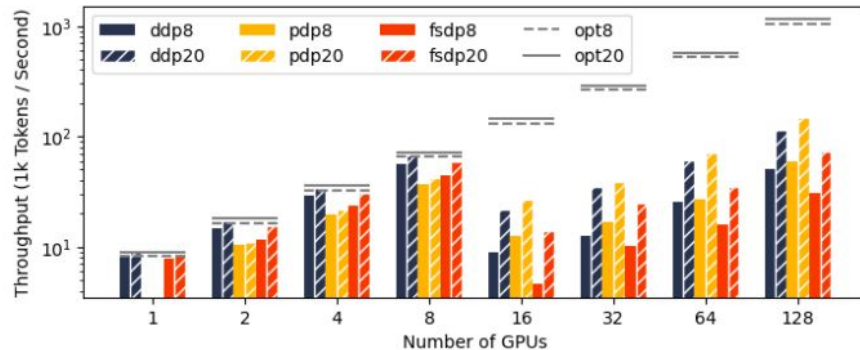


Figure 5. GPT2.7B Throughput vs Number of GPUs




Figure 4. GPTLarge (760M) Throughput vs Number of GPUs

For small, medium sized models, FSDP suffers severely from Communication Overhead.

5.5. Performance Analysis




Figure 10. Throughput vs Model Size (162M — 2.8B)




Figure 11. Peak GPU Memory vs Model Size (162M — 2.8B)

**But FSDP can train large models,
with low memory footprint.**

5.5. Performance Analysis




Figure 12. Throughput vs Model Size (6.8B — 175B)

Only FSDP can train extremely large models.




Figure 13. Peak GPU Memory vs Model Size (6.8B — 175B)

5.5. Performance Analysis




Figure 20. Paradigm Recommendations for 100Gbps Ethernet

Acknowledgement

PYTORCH DISTRIBUTED OVERVIEW [https://pytorch.org/tutorials/beginner/dist_overview.html]

GETTING STARTED WITH DISTRIBUTED DATA PARALLEL [https://pytorch.org/tutorials/intermediate/ddp_tutorial.html]

WRITING DISTRIBUTED APPLICATIONS WITH PYTORCH [https://pytorch.org/tutorials/intermediate/dist_tuto.html]

DISTRIBUTED RPC FRAMEWORK [<https://pytorch.org/docs/stable/rpc.html>]

Launching and configuring distributed data parallel applications
[<https://github.com/pytorch/examples/blob/main/distributed/ddp/README.md>]

PIPELINE PARALLELISM [<https://pytorch.org/docs/stable/pipeline.html>]

TRAINING TRANSFORMER MODELS USING DISTRIBUTED DATA PARALLEL AND PIPELINE PARALLELISM
[https://pytorch.org/tutorials/advanced/ddp_pipeline.html]

ELASTIC AGENT [<https://pytorch.org/docs/stable/elastic/agent.html>]

Acknowledgement

All-Reduce Implementation approaches [<https://algopoolja.tistory.com/95>]

Automatic Cross-Replica Sharding of Weight Update in Data-Parallel Training [<https://arxiv.org/pdf/2004.13336.pdf>]

Fully Sharded Data Parallel: faster AI training with fewer GPUs [<https://engineering.fb.com/2021/07/15/open-source/fsdp/>]

Paradigms of Parallelism [https://colossalai.org/docs/concepts/paradigms_of_parallelism/]

GPipe: Easy Scaling with Micro-Batch Pipeline Parallelism [<https://arxiv.org/pdf/1811.06965.pdf>]

Parallelism in Distributed Deep Learning [<https://insujang.github.io/2022-06-11/parallelism-in-distributed-deep-learning/>]

Zero Offloading brief explanation

[<https://moon-walker.medium.com/large-model-%ED%95%99%EC%8A%B5%EC%9D%98-game-changer-ms%EC%9D%98-deepspeed-zero-1-2-3-%EA%B7%B8%EB%A6%AC%EA%B3%A0-zero-infinity-74c9640190de>]

ZeRO-Offload: Democratizing Billion-Scale Model Training [<https://arxiv.org/df/2101.06840.pdf>]

PyTorch Data Parallel Best Practices on Google Cloud [<https://medium.com/pytorch/pytorch-data-parallel-best-practices-on-google-cloud-6c8da2be180d>]