

Profiling and analyzing GraphSaint with MH-Aug

20230503
박현우

Algorithm 1 GraphSAINT using MH-Aug

Input: target distribution P , proposal distribution Q , original graph G

Output: GCN parameter θ

Initialize: $t \leftarrow 0, G_s^{(0)} \leftarrow G_s$

```
1: while not convergence do
2:   while True do
3:     Draw  $G'$  from  $Q(G'|G^{(t)})$ 
4:     Draw  $u$  from  $Uniform(0, 1)$ 
5:      $G^{(t+1)} \leftarrow G'$ 
6:     if  $u < A$  then
7:       break
8:     end if
9:   end while
10:  for each batch do
11:     $G_s(V_s, E_s) \leftarrow$  Sampled subgraph of  $G$ 
12:     $G_s^{(t)}(V_s, E_s) \leftarrow$  Sampled subgraph of  $G^{(t)}$ 
13:     $G_s^{(t+1)}(V_s, E_s) \leftarrow$  Sampled subgraph of  $G^{(t+1)}$ 
14:    Update  $\theta$  with  $L(G_s, G_s^{(t)}, G_s^{(t+1)}, \theta)$ 
15:     $t \leftarrow t + 1$ 
16:  end for
17: end while
```

← Augmentation Part

10/88

Training is very slow and as a result,
the model doesn't converge well.

= Are there any bottlenecks?

Run profiling in training of one
sample graph in the dataset.

```

def generate_aug_graph(g, model,
                      sigma_delta_e=0.03, sigma_delta_v=0.03, mu_e=0.6, mu_v=0.2,
                      lam1_e=1, lam1_v=1, lam2_e=0.0, lam2_v=0.0,
                      a_e=100, b_e=1, a_v=100, b_v=1):
    # Original Graph Feature and Metadata Extraction, Preprocessing
    num_nodes = g.num_nodes()
    num_edges = g.num_edges()

    coo_mat = g.edges(form='uv')
    coo_mat = th.tensor([list(coo_mat[0]), list(coo_mat[1])], device='cuda:0')
    n_list = th.ones(num_nodes)

    # Create Aggregate Model
    agg_model = AGGNet(num_hop=2)
    agg_model.cuda()

    while True:
        # Calculate Delta Value
        delta_G_e = 1 - coo_mat.shape[1] / num_edges
        delta_G_e_aug = our_truncnorm(0, 1, delta_G_e, sigma_delta_e, mode='rvs')

        delta_G_v = 1 - n_list.sum().item() / num_nodes
        delta_G_v_aug = our_truncnorm(0, 1, delta_G_v, sigma_delta_v, mode='rvs')

        # Graph Augmentation According To Delta Value
        aug_g, aug_n_list = augment(g, delta_G_e_aug, delta_G_v_aug)
        aug_g = dgl.add_self_loop(aug_g)

        # message_passing_g = copy.deepcopy(g)
        message_passing_g = g.clone()
        message_passing_g.ndata['feat'] = th.ones(num_nodes, 1, device='cuda:0')

```

```

# Split into sub functions for profiling.
def _generate_aug_graph(g, model,
                       sigma_delta_e=0.03, sigma_delta_v=0.03, mu_e=0.6, mu_v=0.2,
                       lam1_e=1, lam1_v=1, lam2_e=0.0, lam2_v=0.0,
                       a_e=100, b_e=1, a_v=100, b_v=1):

    # Original Graph Feature and Metadata Extraction, Preprocessing
    def initialize():
        num_nodes = g.num_nodes()
        num_edges = g.num_edges()

        coo_mat = g.edges(form='uv')
        coo_mat = torch.tensor([list(coo_mat[0]), list(coo_mat[1])]).to(DEVICE)
        n_list = torch.ones(num_nodes).to(DEVICE)

        # Create Aggregate Model
        agg_model = AGGNet(num_hop=2)
        agg_model.cuda()

        return num_nodes, num_edges, coo_mat, n_list, agg_model

    def calculate_delta_value():
        # Calculate Delta Value
        delta_G_e = 1 - coo_mat.shape[1] / num_edges
        delta_G_e_aug = our_truncnorm(0, 1, delta_G_e, sigma_delta_e, mode='rvs')

        delta_G_v = 1 - n_list.sum().item() / num_nodes
        delta_G_v_aug = our_truncnorm(0, 1, delta_G_v, sigma_delta_v, mode='rvs')

        return delta_G_e, delta_G_e_aug, delta_G_v, delta_G_v_aug

```

Algorithm 1 GraphSAINT using MH-Aug

Input: target distribution P , proposal distribution Q , original graph G

Output: GCN parameter θ

Initialize: $t \leftarrow 0, G_s^{(0)} \leftarrow G_s$

```
1: while not convergence do
2:   while True do
3:     Draw  $G'$  from  $Q(G'|G^{(t)})$ 
4:     Draw  $u$  from  $Uniform(0,1)$ 
5:      $G^{(t+1)} \leftarrow G'$ 
6:     if  $u < A$  then ← Augmentation Part
7:       break
8:     end if
9:   end while
10:  for each batch do
11:     $G_s(V_s, E_s) \leftarrow$  Sampled subgraph of  $G$ 
12:     $G_s^{(t)}(V_s, E_s) \leftarrow$  Sampled subgraph of  $G^{(t)}$ 
13:     $G_s^{(t+1)}(V_s, E_s) \leftarrow$  Sampled subgraph of  $G^{(t+1)}$ 
14:    Update  $\theta$  with  $L(G_s, G_s^{(t)}, G_s^{(t+1)}, \theta)$ 
15:     $t \leftarrow t + 1$ 
16:  end for
17: end while
```

10/88

```
with record_function("initialize"):
    num_nodes, num_edges, coo_mat, n_list, agg_model = initialize()

while True:

    with record_function("calculate_delta_value"):
        delta_g_e, delta_g_e_aug, delta_g_v, delta_g_v_aug = calculate_delta_value()

    with record_function("graph_augmentation"):
        aug_g = graph_augmentation()

    with record_function("message_passing"):
        message_passing_g, message_passing_aug_g = message_passing()

    with record_function("calculate_ego_graph_message_passing_value"):
        org_ego = calculate_ego_graph_message_passing_value()

    with record_function("calculate_augmented_delta"):
        delta_g_e, delta_g_e_aug, delta_g_v, delta_g_v_aug = calculate_augmented_delta()

    with record_function("calculate_distribution"):
        h_loss_op = calculate_distribution()

    with record_function("compute_model"):
        output = compute_model()

    with record_function("compute_ent"):
        ent = compute_ent()

    with record_function("compute_p_aug"):
        p, p_aug = compute_p_aug()

    with record_function("compute_q_aug"):
        q, q_aug = compute_q_aug()
```

Configuration

Number of Worker(s)


4

Device Type


CPU

Execution Summary

Category	Time Duration (us)	Percentage (%)
Average Step Time	6,957,929	100
DataLoader	0	0
CPU Exec	5,294,977	76.1
Other	1,662,952	23.9




Step Time Breakdown ②




Operator View

☐ All operators ☒ Top operators to show 10

Host Self Time (us) ⓘ



Host Total Time (us) ⓘ



(left) Host Self Time

: Accumulated time of operator itself only.


(right) Host Total Time

: Accumulated time of both operator and its all children operators.


Operator View

☐ All operators ☒ Top operators to show 10

Host Self Time (us)



Host Total Time (us)



(left) / (right 2)

aten::_local_scalar_dense

Dense operation


(right) aten::item

Ordered_dict for
parameter_dict,
containing and handling
parameters.

Operator View


☐ All operators ☒ Top operators to show 10

Host Self Time (us) ?



aten::_local_scalar_dense
aten::copy_
aten::mm
GSplitM
aten::item
aten::select
aten::unbind
aten::_to_copy
aten::as_strided
aten::mul

Host Total Time (us) ?



aten::item
aten::_local_scalar_dense
aten::to
aten::_to_copy
aten::copy_
aten::matmul
aten::mm
GSplitM
aten::unbind
aten::select

(left) aten::copy_

Copy from device
memory into host
memory

(right) aten::to


Move memory to other
memory (H2D or D2H)

(right) aten::_to_copy

Copy from host memory
into device memory

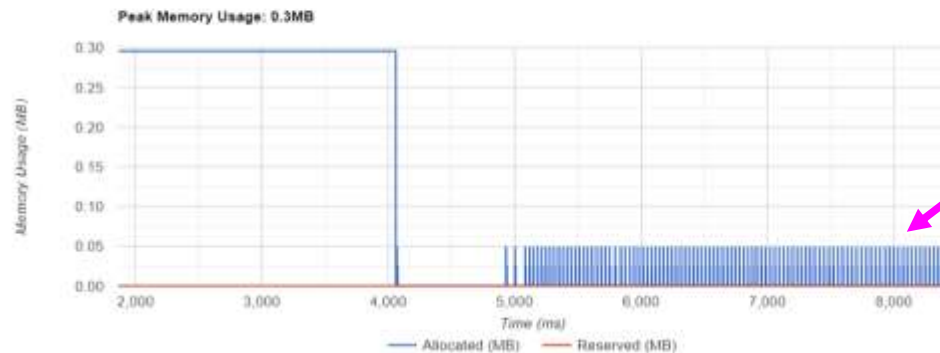
Memory View

Device:
GPU0



Memory View

Device:
CPU



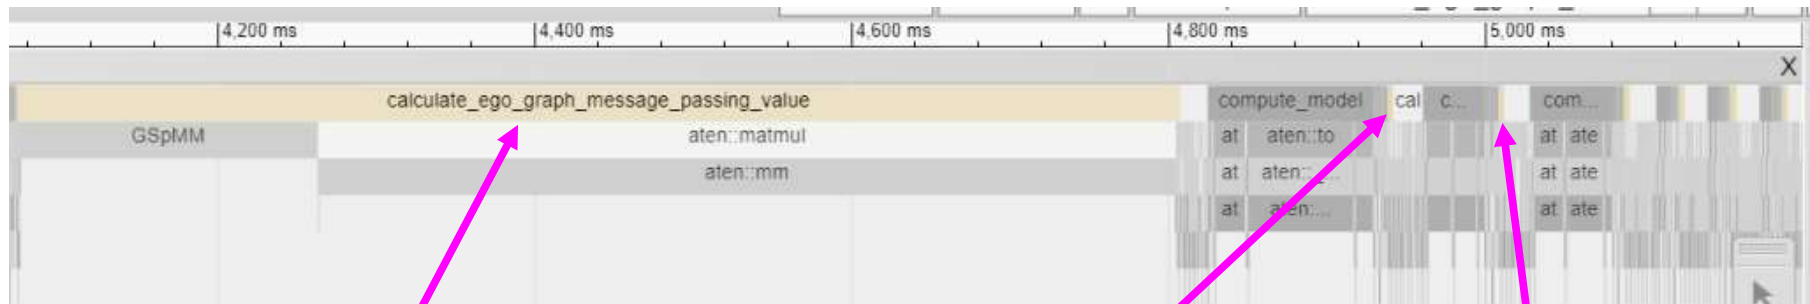
A lot of H2D and
D2H copies



calculate_ego_graph_message_
passing_value

Initialization

Augmentation loop



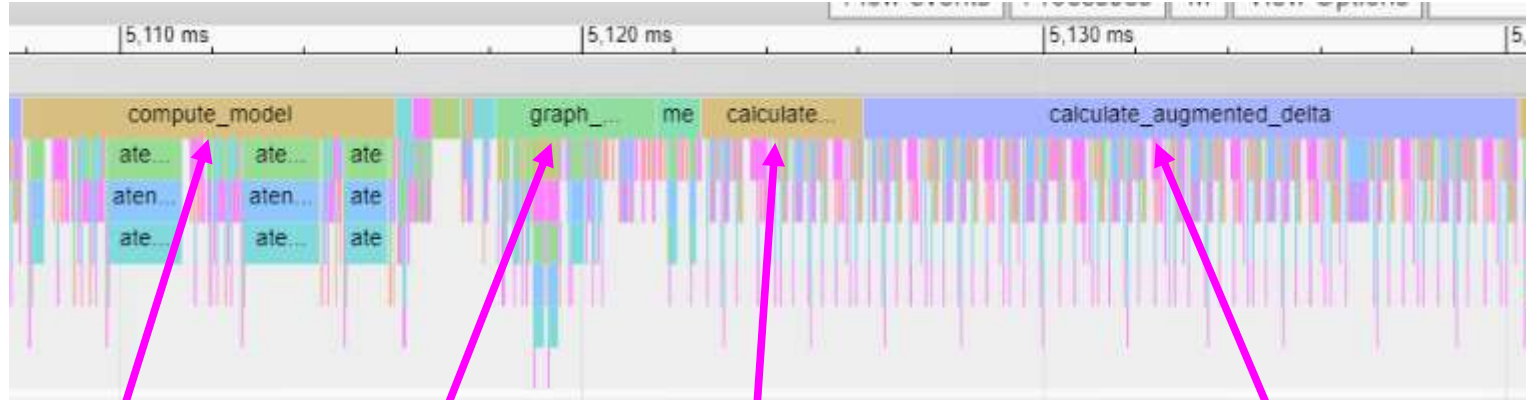
1 item selected.		Slice (1)
Title	calculate_ego_graph_message_passing_value	
Category	user_annotation	
User Friendly Category	other	
Start	4,073.400 ms	
Wall Duration	734.477 ms	
Self Time	2.563 ms	
▼Args		
External id	155217	
Ev Idx	155216	

1 item selected.		Slice (1)
Title	calculate_ego_graph_message_passing_value	
Category	user_annotation	
User Friendly Category	other	
Start	4,938.216 ms	
Wall Duration	3.814 ms	
Self Time	1.772 ms	
▼Args		
External id	156306	
Ev Idx	156305	

1 item selected.		Slice (1)
Title	calculate_ego_graph_message_passing_value	
Category	user_annotation	
User Friendly Category	other	
Start	5,008.656 ms	
Wall Duration	3.586 ms	
Self Time	1.834 ms	
▼Args		
External id	157395	
Ev Idx	157394	

May be warming up? (= loaded on device and cached)

One cycle of augmentation loop



1. `compute_model`

2. `graph_augmentation`

3. `calculate_ego_graph_message_passing_value`

4. `calculate_augmented_delta`

Top 4 Time taken : $4 > 1 > 3 > 2$

```
def calculate_ego_graph_message_passing_value():
    # Calculate ego-graph's message passing value
    with torch.no_grad():
        org_ego = aggregate(message_passing_g, agg_model)

    return org_ego
```

```
def compute_model():
    with torch.no_grad():
        output = model(g)

    return output
```

```
def calculate_augmented_delta():
    # Calculate Augmented Delta Value
    with torch.no_grad():
        delta_g_e = 1 - (aggregate(message_passing_g, agg_model) / org_ego).squeeze(1)
        delta_g_aug_e = 1 - (aggregate(message_passing_aug_g, agg_model) / org_ego).squeeze(1)
        delta_g_v = 1 - (aggregate(message_passing_g, agg_model) / org_ego).squeeze(1)
        delta_g_aug_v = 1 - (aggregate(message_passing_aug_g, agg_model) / org_ego).squeeze(1)

    return delta_g_e, delta_g_aug_e, delta_g_v, delta_g_aug_v
```

```
def graph_augmentation():
    # Graph Augmentation According To Delta Value
    aug_g, aug_n_list = augment(g, delta_G_e_aug, delta_G_v_aug)
    aug_g = dgl.add_self_loop(aug_g)

    return aug_g
```

They use aggregate()
function
= aggregate() is the
bottleneck?

```

class AGGNet(nn.Module):
    def __init__(self, num_hop, in_feats=1, hidden_feats=1, out_feats=1, dropout=0, use_bn=False):
        super(AGGNet, self).__init__()
        self.use_bn = use_bn
        self.dropout = dropout
        self.num_hop = num_hop

        self.convs = nn.ModuleList()

        if self.use_bn:
            self.bns = nn.ModuleList()
            self.bns.append(nn.BatchNorm1d(hidden_feats))

        self.convs.append(GraphConv(in_feats, out_feats, allow_zero_in_degree=True))

        for _ in range(self.num_hop - 1):
            if self.use_bn:
                self.bns.append(nn.BatchNorm1d(hidden_feats))
            self.convs.append(GraphConv(out_feats, out_feats, allow_zero_in_degree=True))

    def forward(self, graph):
        for i in range(self.num_hop - 1):
            feat = graph.ndata['feat']

            # graph = dgl.add_self_loop(graph)
            feat = self.convs[i](graph, feat)
            if self.use_bn:
                feat = self.bns[i](feat)
            feat = F.dropout(feat, p=self.dropout, training=self.training)
        feat = self.convs[-1](graph, feat)

        return feat

```

```
# from aug.py
```

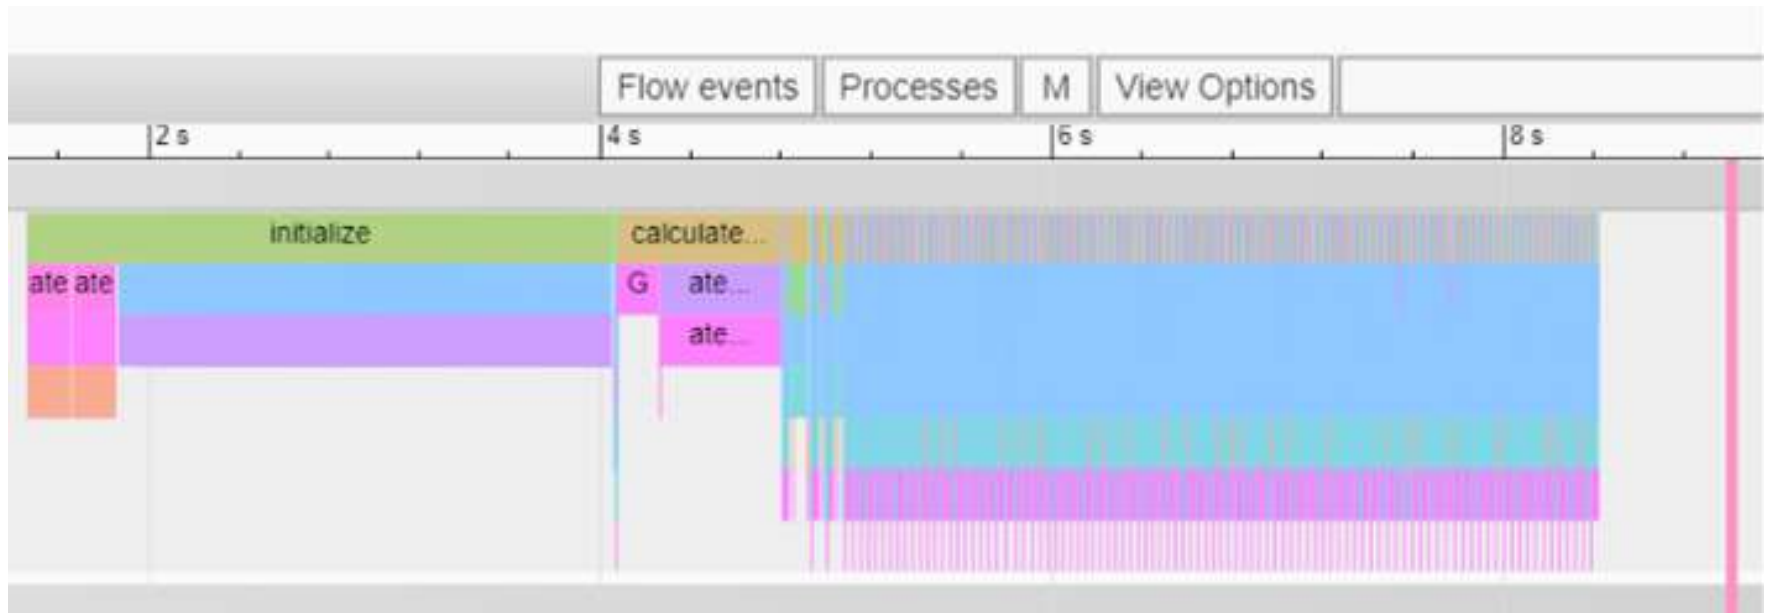
```

def aggregate(graph, agg_model):
    s_vec = agg_model(graph)
    return s_vec

```

Future work #1

Optimize aggregate()
function (= AGGNet)




Augmentation loop

= Inside the loop is repeated a lot of times

TorchDynamo (Torch Script)

TorchDynamo is a Python-level JIT compiler for accelerating PyTorch code.



TorchDynamo (Torch Script)

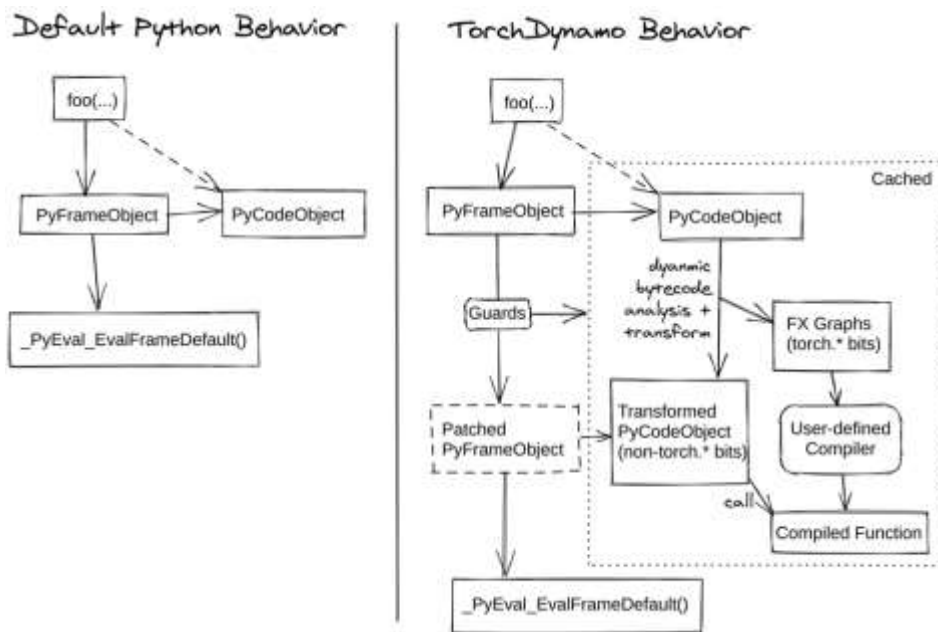
TorchDynamo dynamically modifies or “compiles” Python bytecode right before it is executed.

Then, it rewrites or “optimizes” Python bytecode in order to extract sequences of PyTorch operations.

Extracted PyTorch ops sequences are optimized, thus much faster than plain python codes or normal Pytorch code.

TorchDynamo

Furthermore, once compiled, these extracted PyTorch ops sequences are Cached.



```

with record_function("initialize"):
    num_nodes, num_edges, coo_mat, n_list, agg_model = initialize()

while True:

    with record_function("calculate_delta_value"):
        delta_G_e, delta_G_e_aug, delta_G_v, delta_G_v_aug = calculate_delta_value()

    with record_function("graph_augmentation"):
        aug_g = graph_augmentation()

    with record_function("message_passing"):
        message_passing_g, message_passing_aug_g = message_passing()

    with record_function("calculate_ego_graph_message_passing_value"):
        org_ego = calculate_ego_graph_message_passing_value()

    with record_function("calculate_augmented_delta"):
        delta_g_e, delta_g_aug_e, delta_g_v, delta_g_aug_v = calculate_augmented_delta()

    with record_function("calculate_distribution"):
        h_loss_op = calculate_distribution()

    with record_function("compute_model"):
        output = compute_model()

    with record_function("compute_ent"):
        ent = compute_ent()

    with record_function("compute_p_aug"):
        p, p_aug = compute_p_aug()

    with record_function("compute_q_aug"):
        q, q_aug = compute_q_aug()

```

Future work #2

Compile loop contexts.

NOTE : TorchDynamo is beta;
 unstable and only works for
 Plain python code and PyTorch
 - unable to apply to external
 libraries;
 e.g. DGL or Sci-py